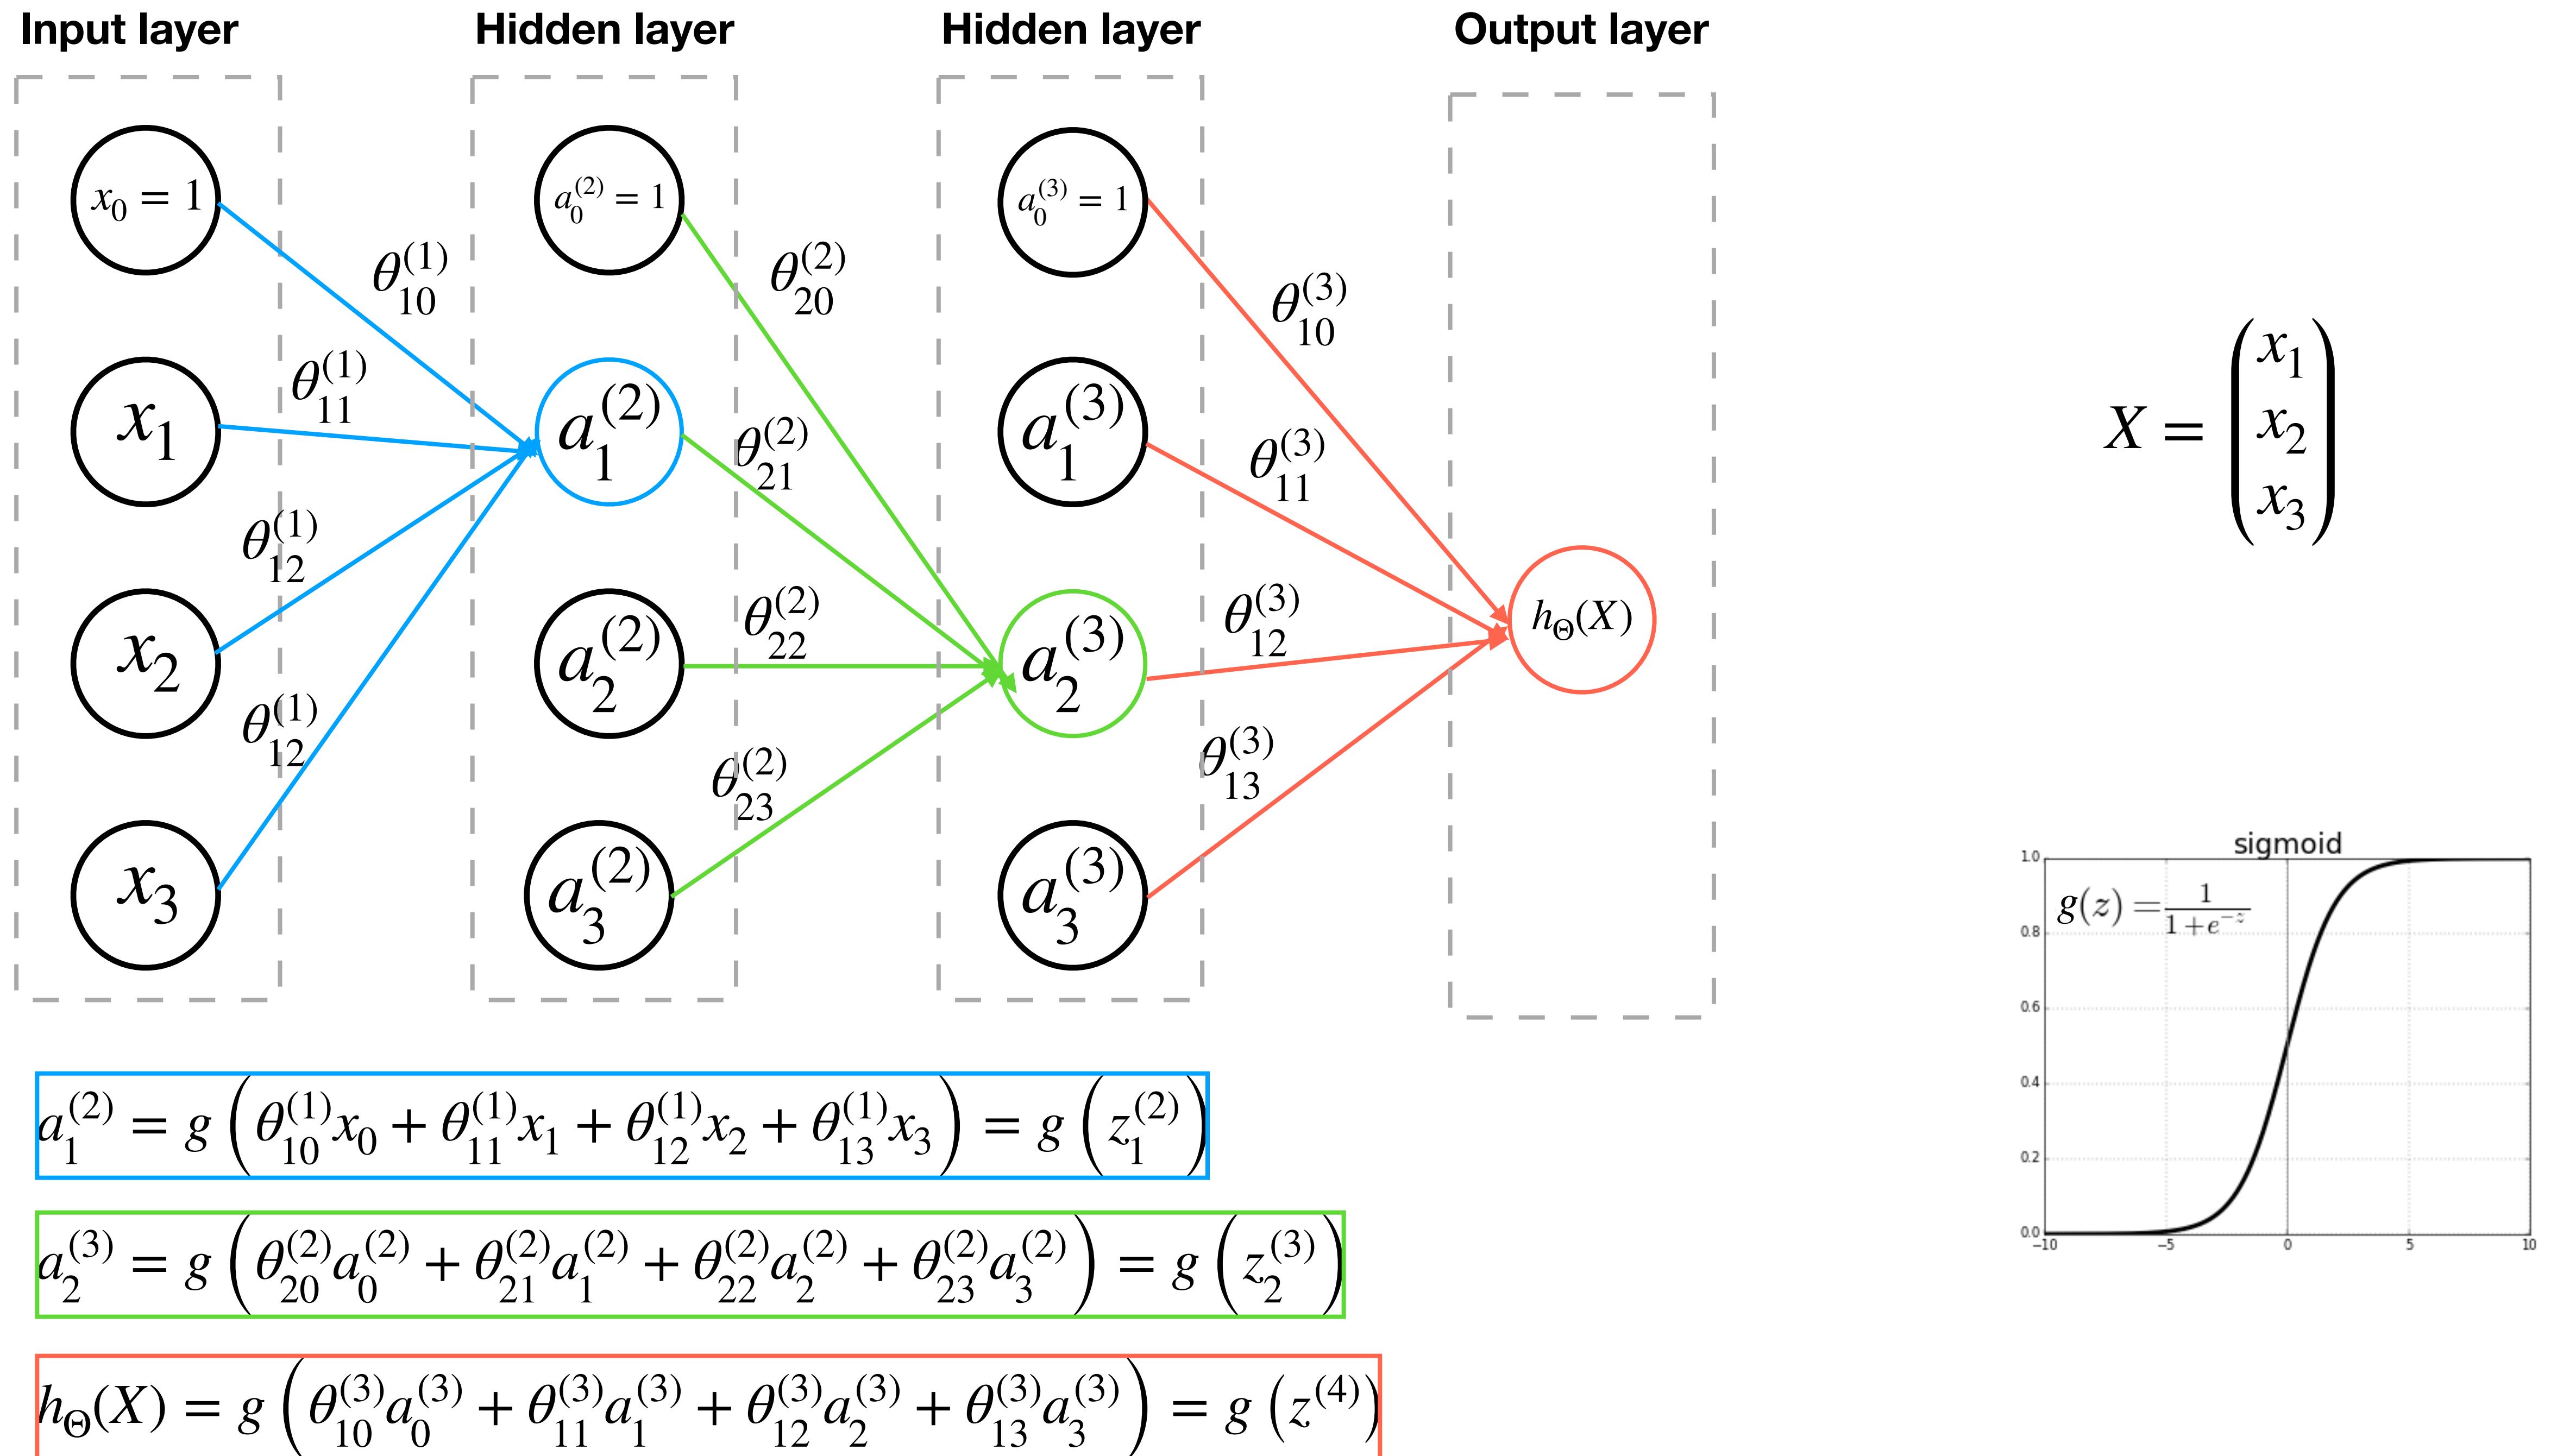


# **Convolutional Neural Networks**

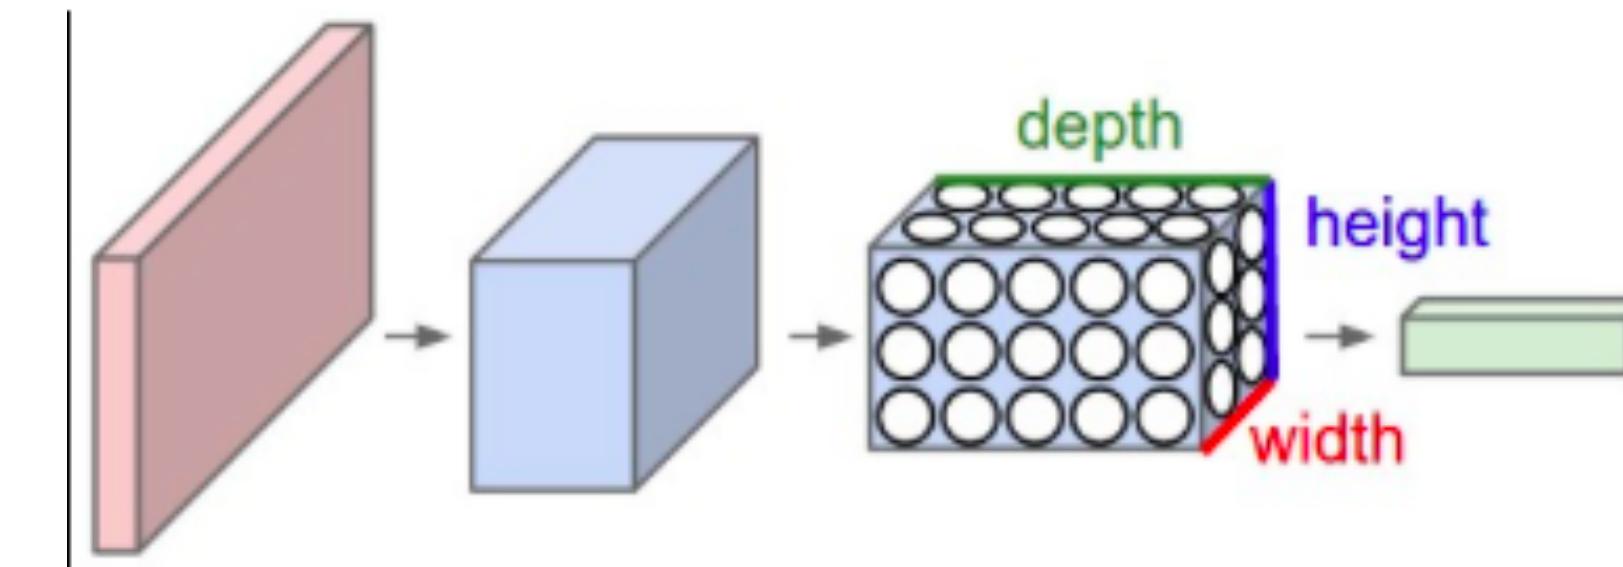
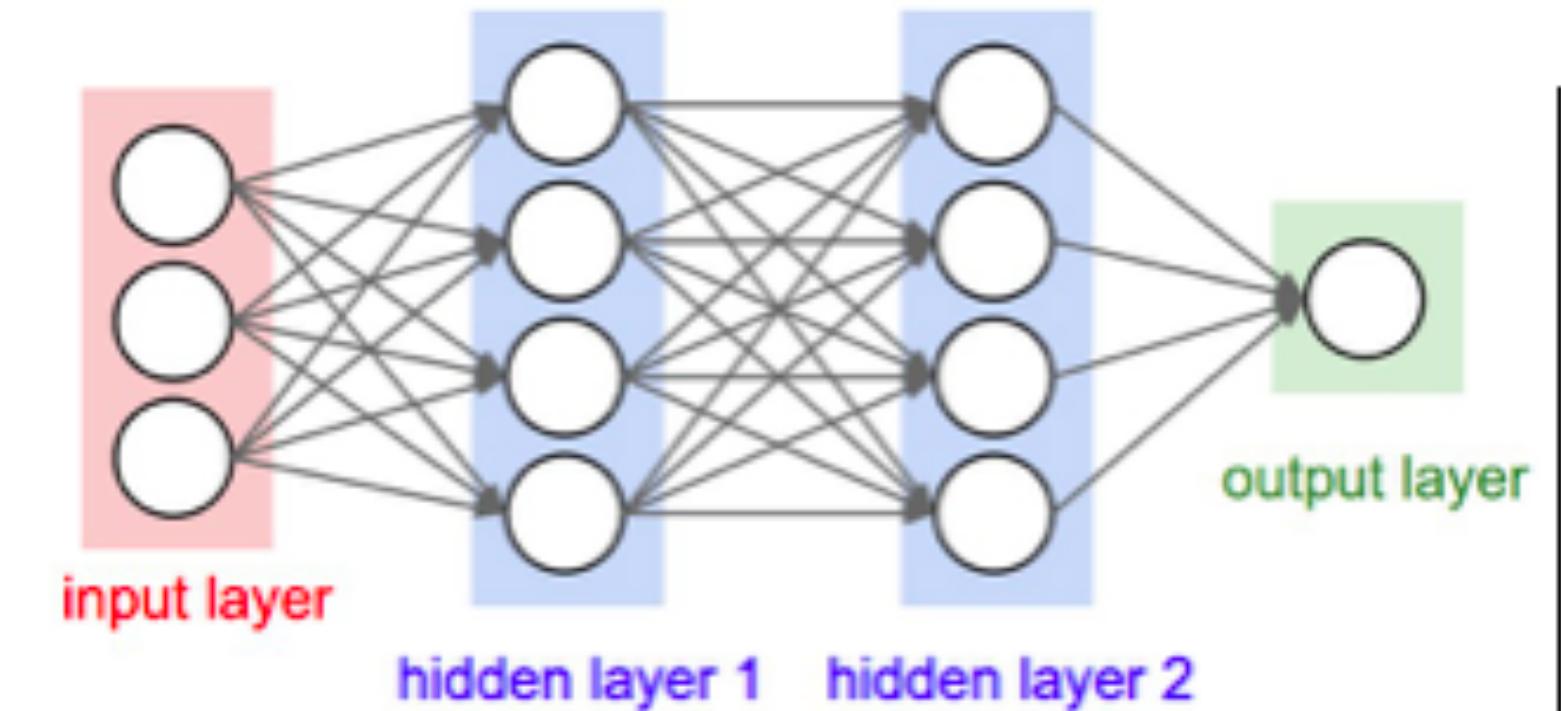
# Convolutional Neural Networks

<https://github.com/JBCA-MachineLearning/Workshop/tree/master/CNN>



# Convolutional layer

- Convolutional layer
- ReLu Layer
- Pooling Layer
- Dropout Layer
- Fully Connected Layer



# Convolution Layer

```
In [ ]: from keras.models import Sequential  
from keras.layers import Conv2D  
  
model = Sequential()  
model.add(Conv2D(32, (3, 3), input_shape=input_shape))
```

**Input**

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

**Filter/Kernel**

1	0	1
0	1	0
1	0	1

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0

4		

# Convolution Layer

```
In [ ]: from keras.models import Sequential  
from keras.layers import Conv2D  
  
model = Sequential()  
model.add(Conv2D(32, (3, 3), input_shape=input_shape))
```

**Input**

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

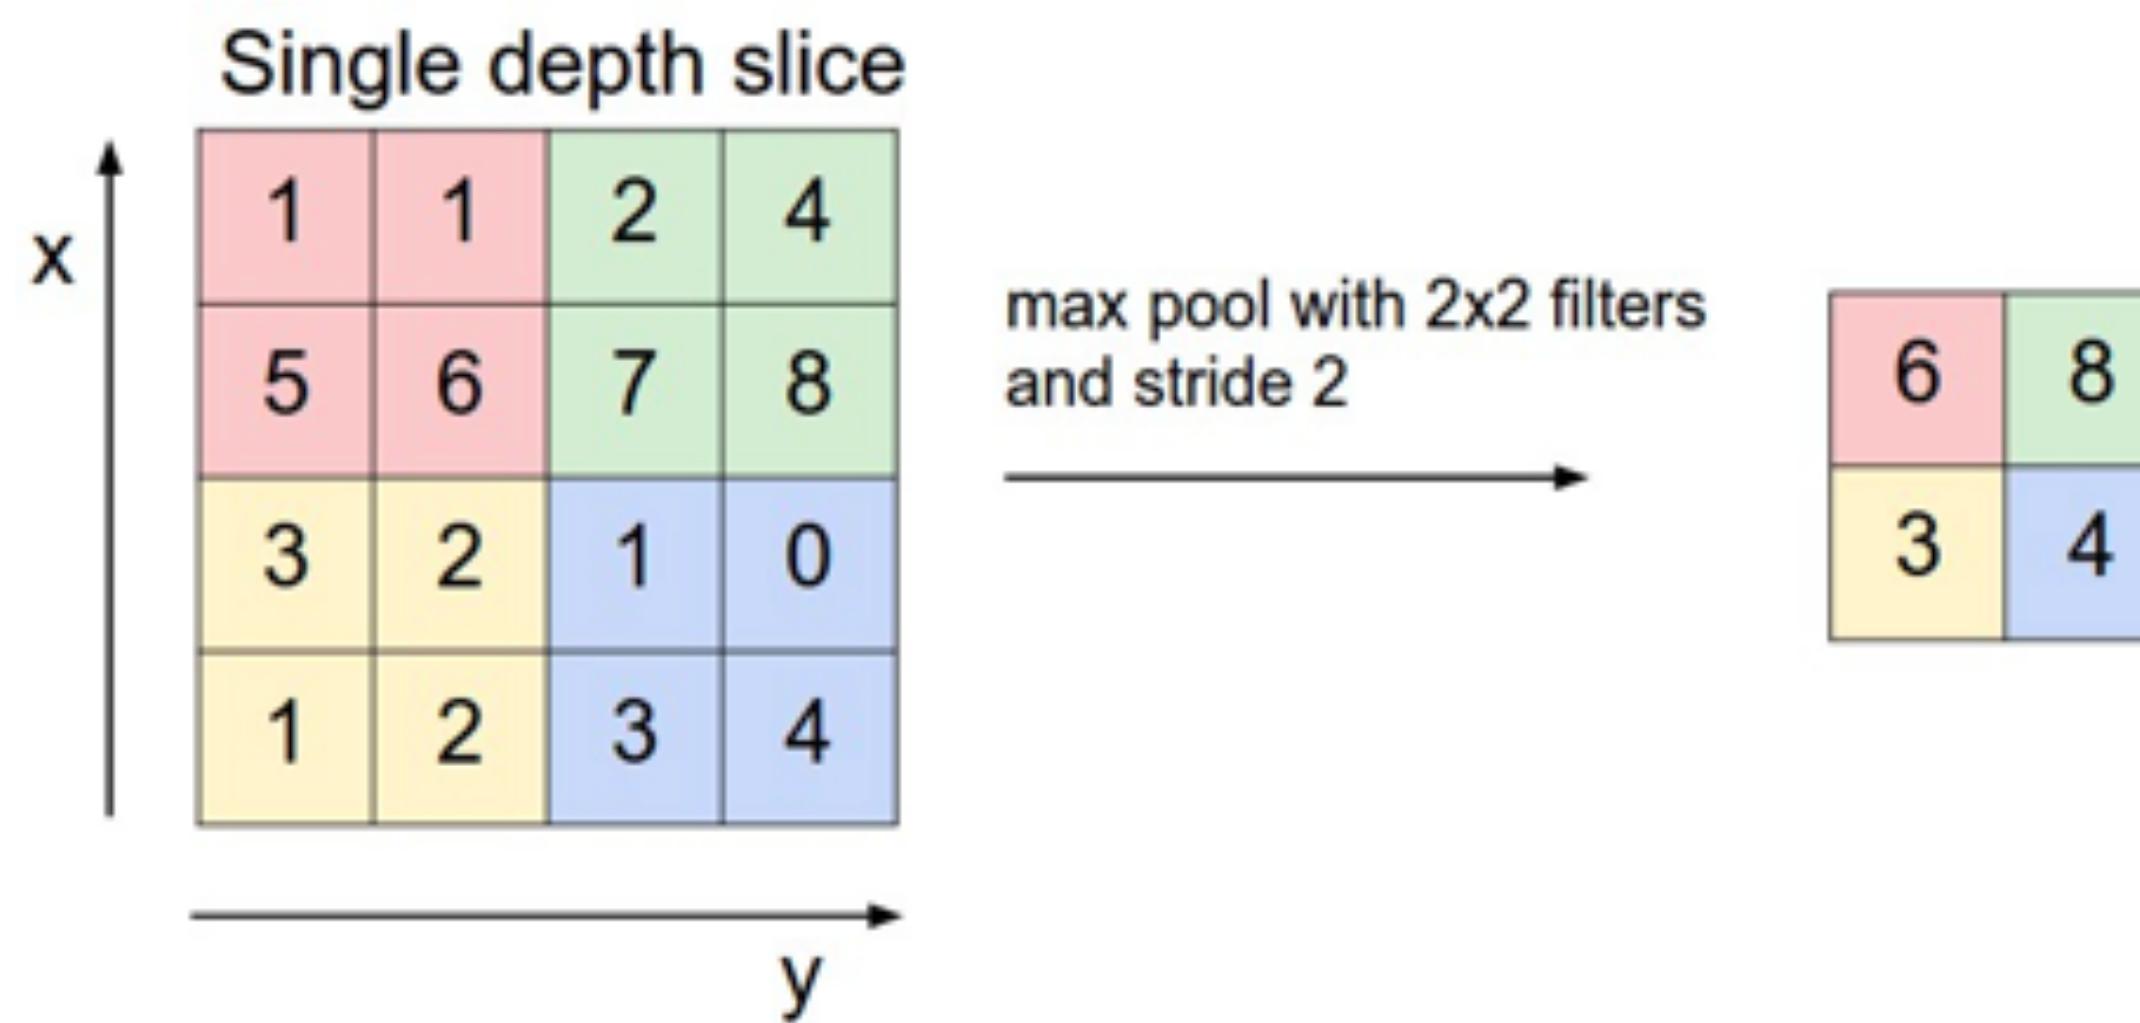
**Filter/Kernel**

1	0	1
0	1	0
1	0	1

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0

4		

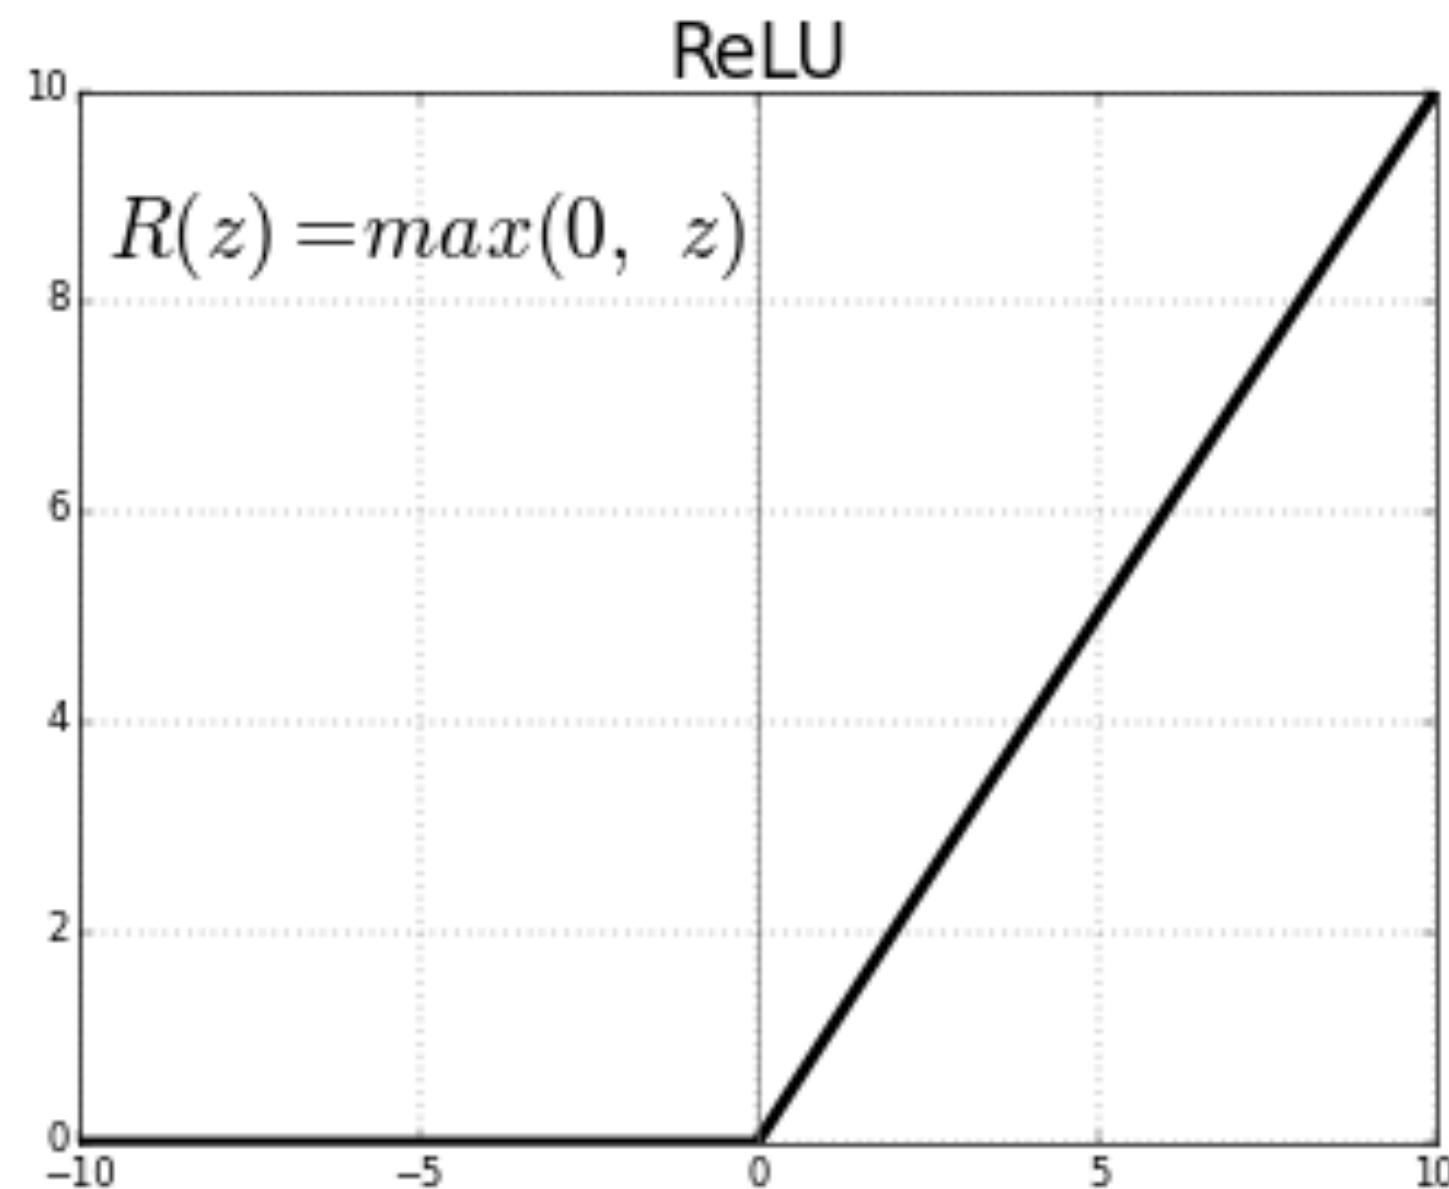
# Pooling layer



- Reduce Image size and therefore number of parameters to be learned
- Can help with overfitting

```
In [ ]: from keras.layers import MaxPooling2D  
model.add(MaxPooling2D(pool_size=(2, 2)))
```

# ReLU layer

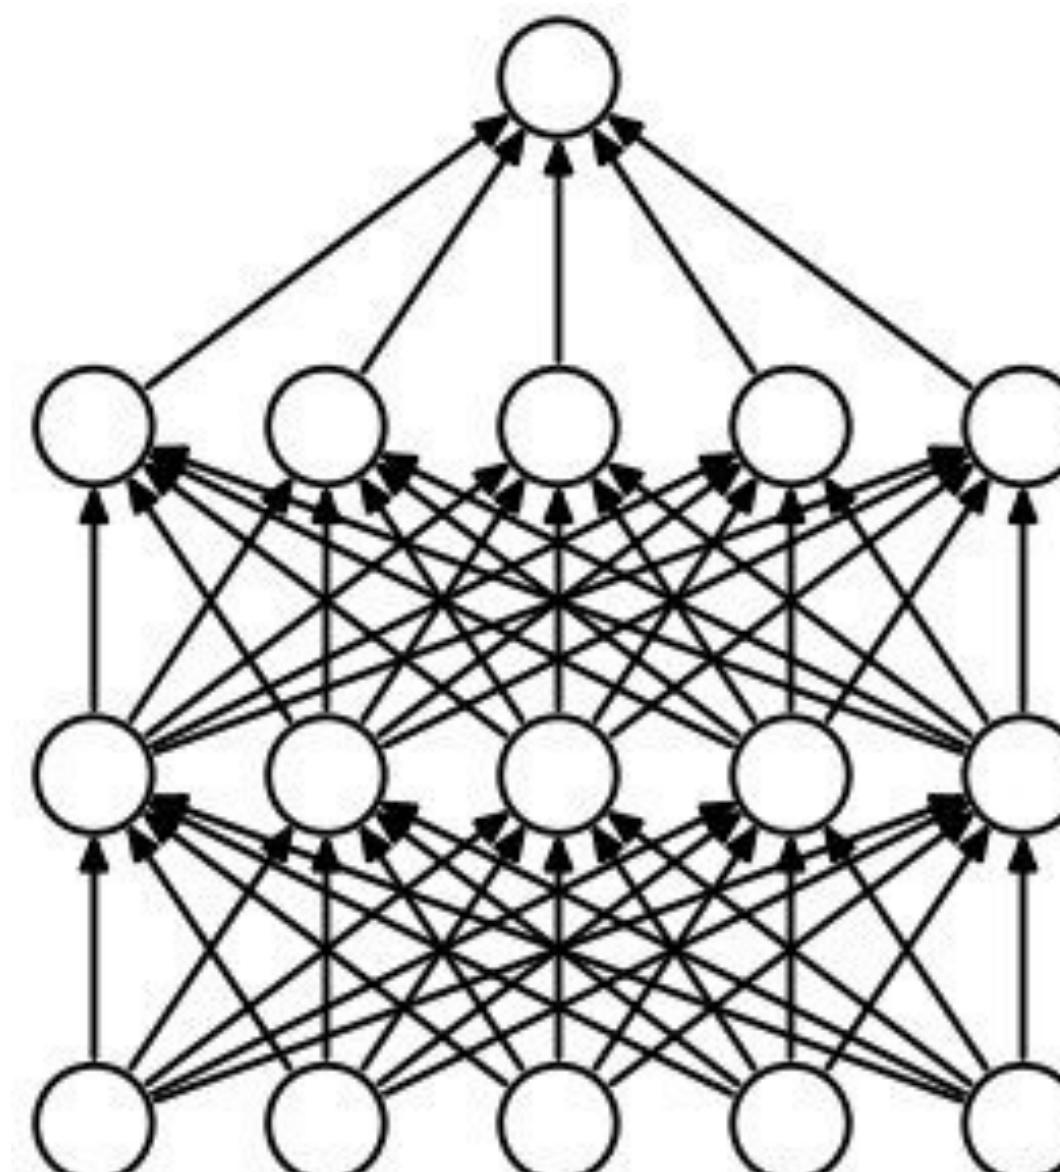


- Added after convolutional layer
- Introduces non-linearity to network

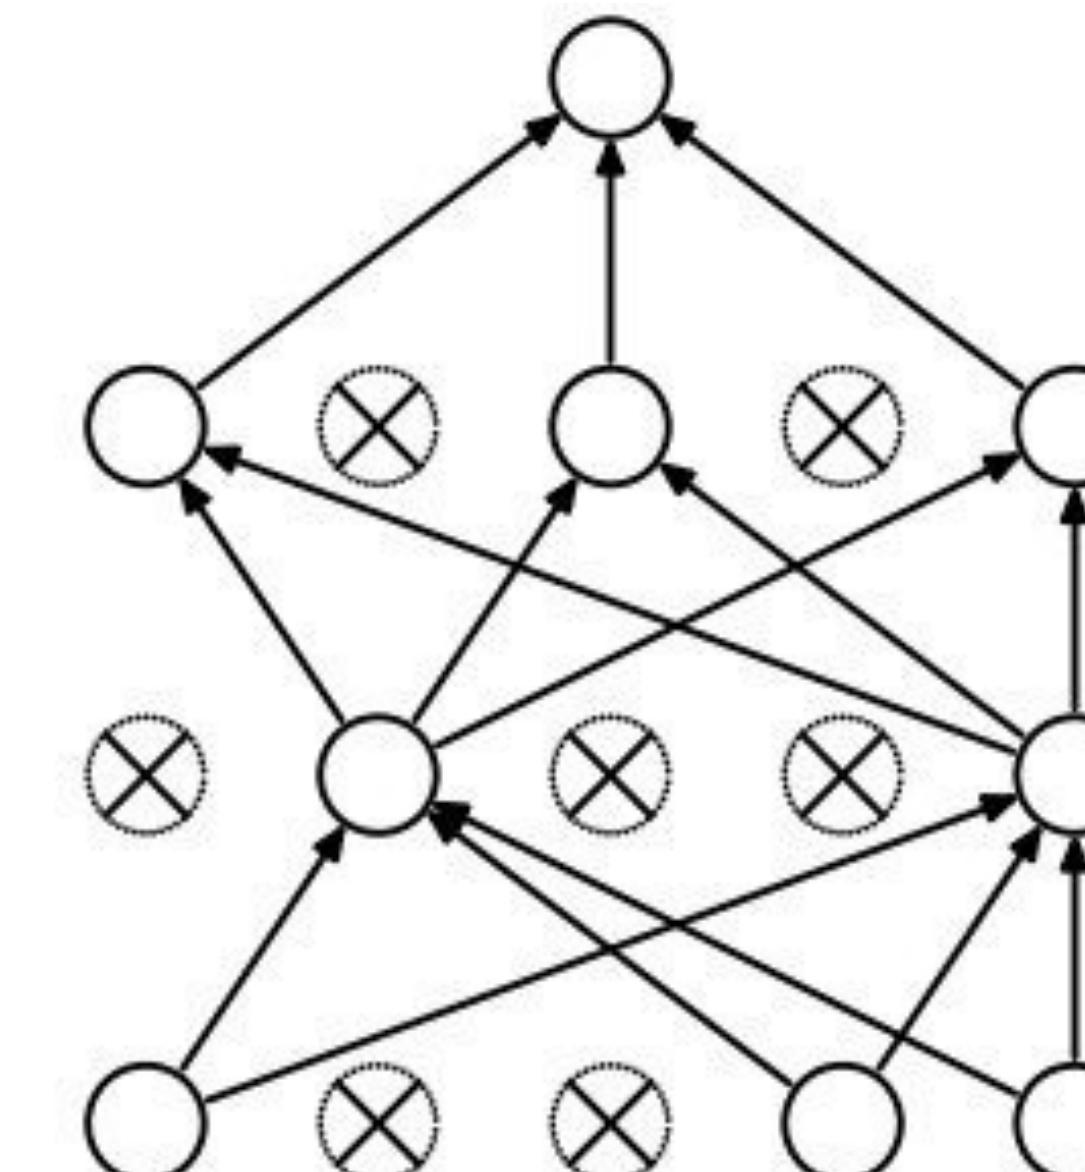
```
In [ ]: from keras.layers import Activation  
model.add(Activation('relu'))
```

# Drop out layer

```
In [ ]: from keras.layers import Activation, Dropout  
model.add(Dropout(0.5))
```



(a) Standard Neural Net



(b) After applying dropout.

# Putting it all together

```
In [2]: from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D
from keras.layers import Activation, Dropout, Flatten, Dense

model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=input_shape))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(64, (3, 3)))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

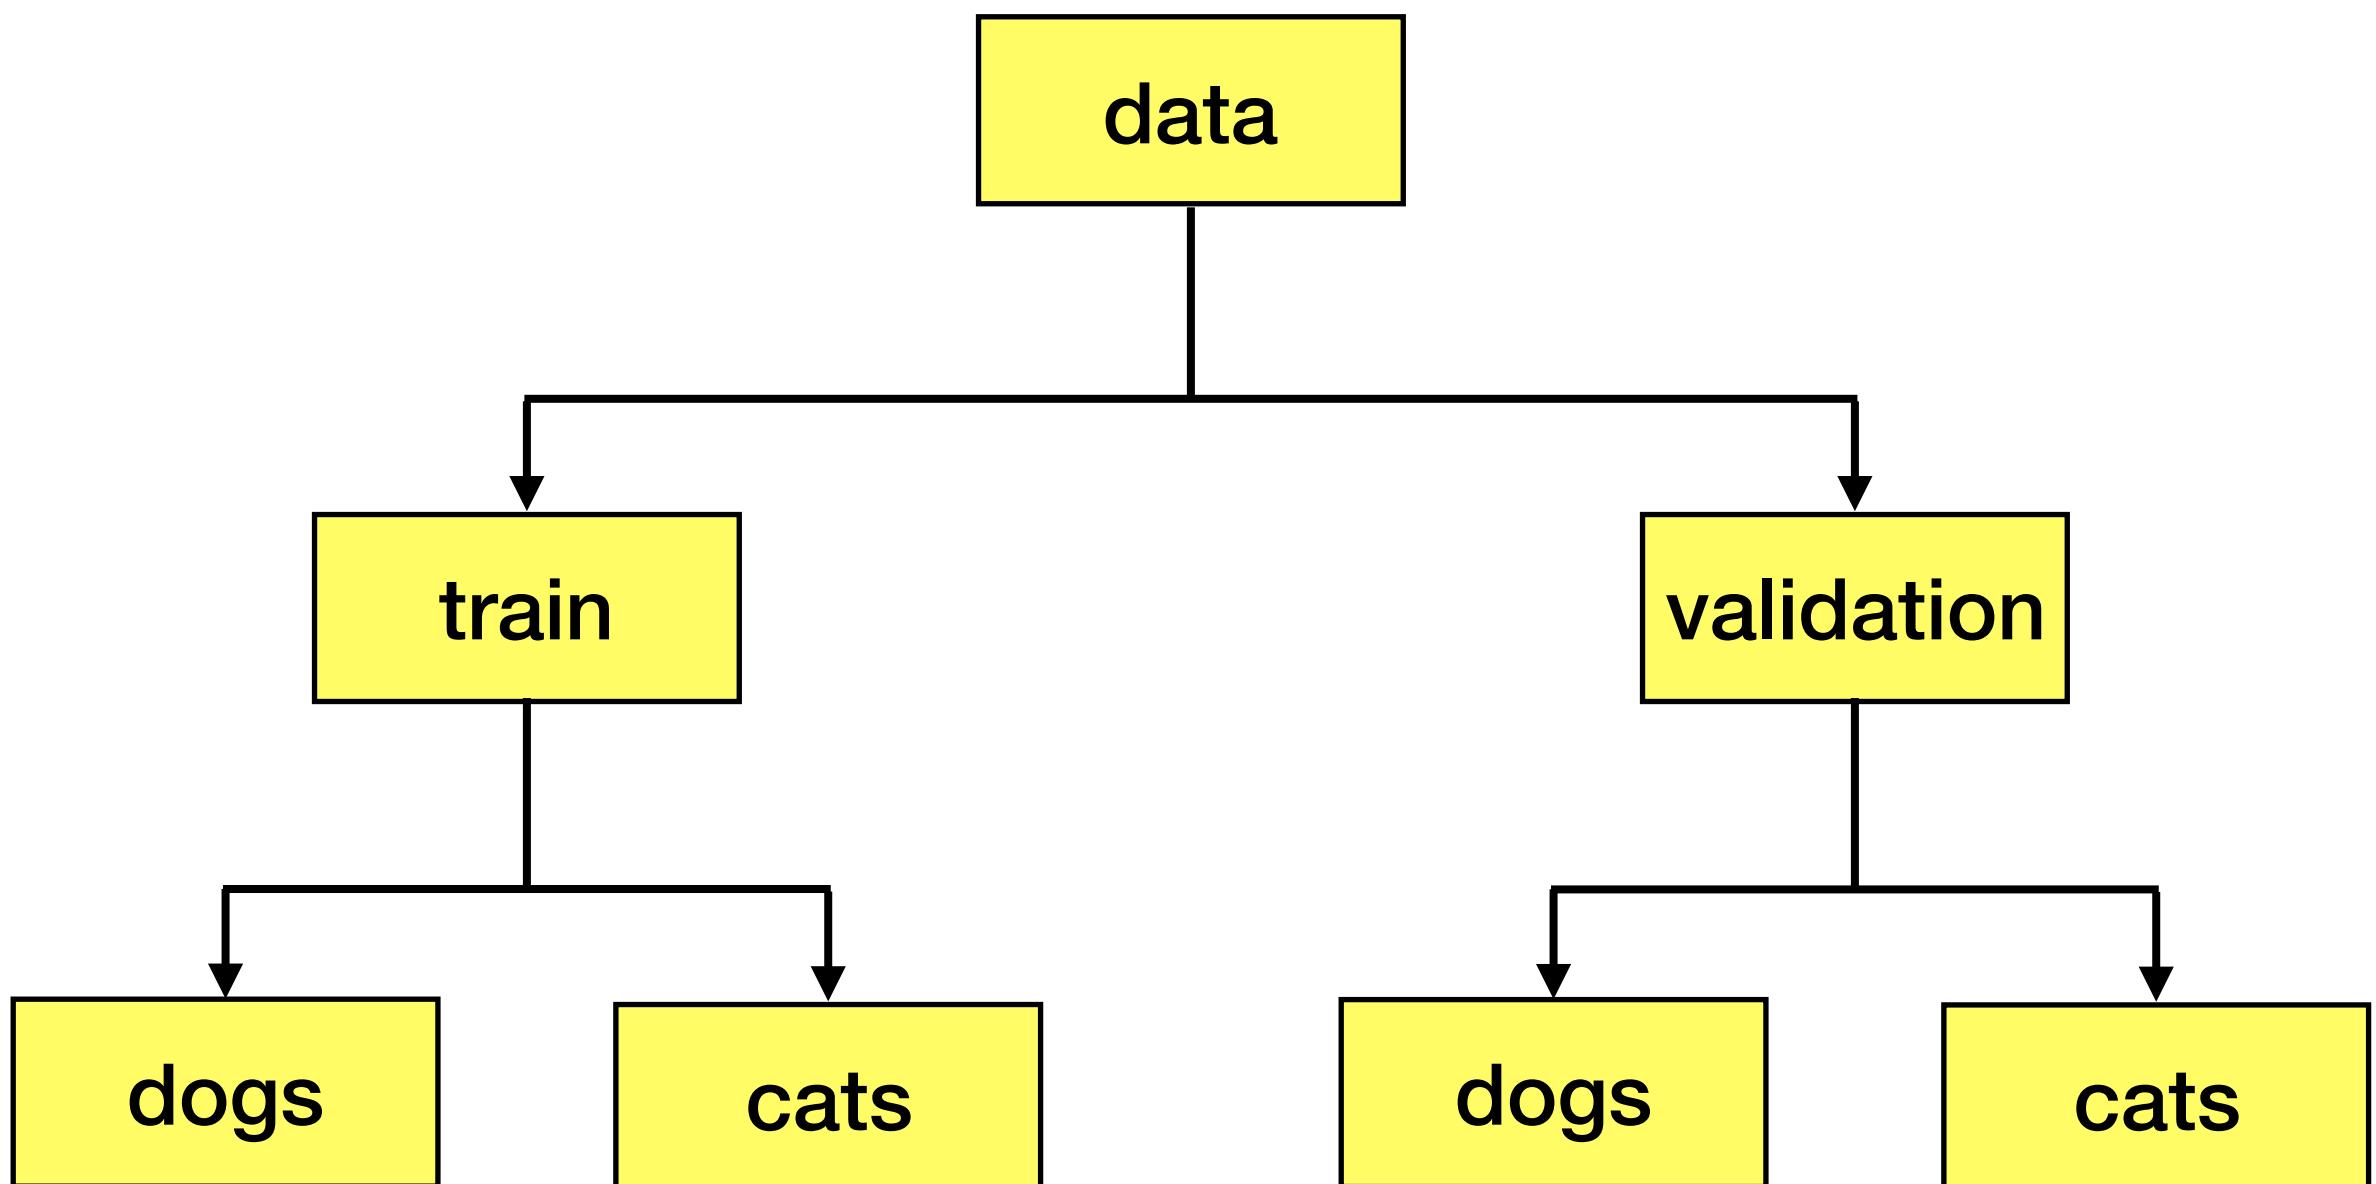
model.add(Flatten())
model.add(Dense(64))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(1))
model.add(Activation('sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

[http://scs.ryerson.ca/~aharley/  
vis/conv/flat.html](http://scs.ryerson.ca/~aharley/vis/conv/flat.html)

Using TensorFlow backend.

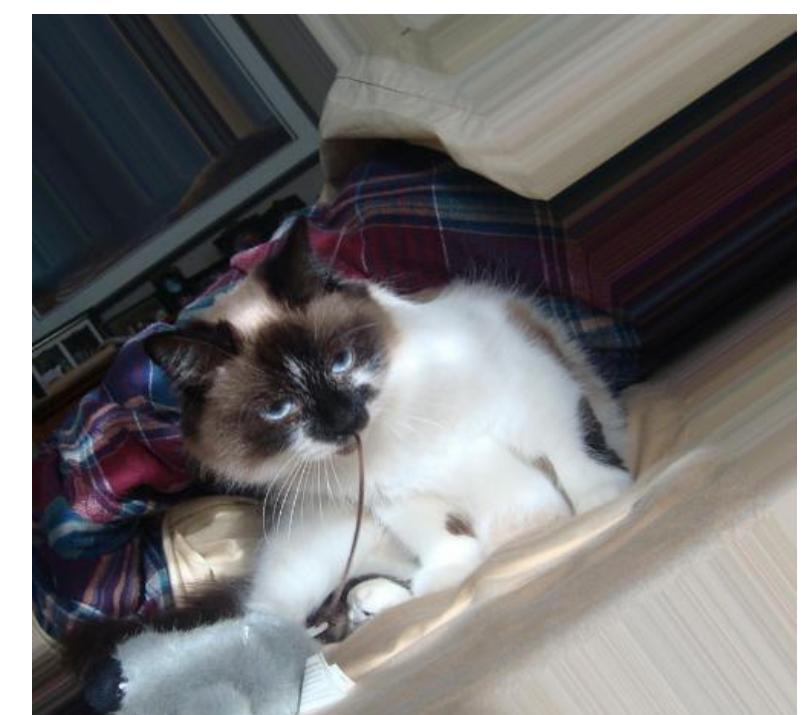
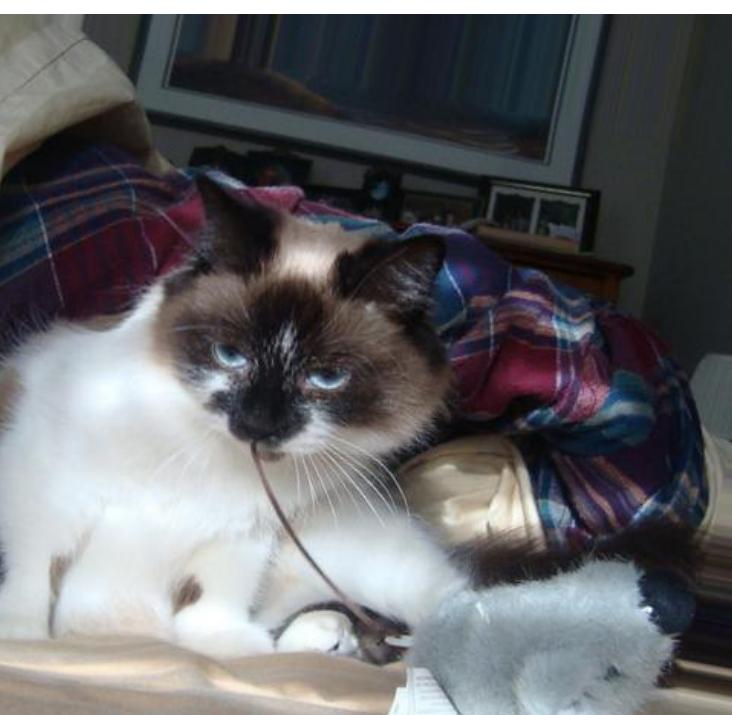
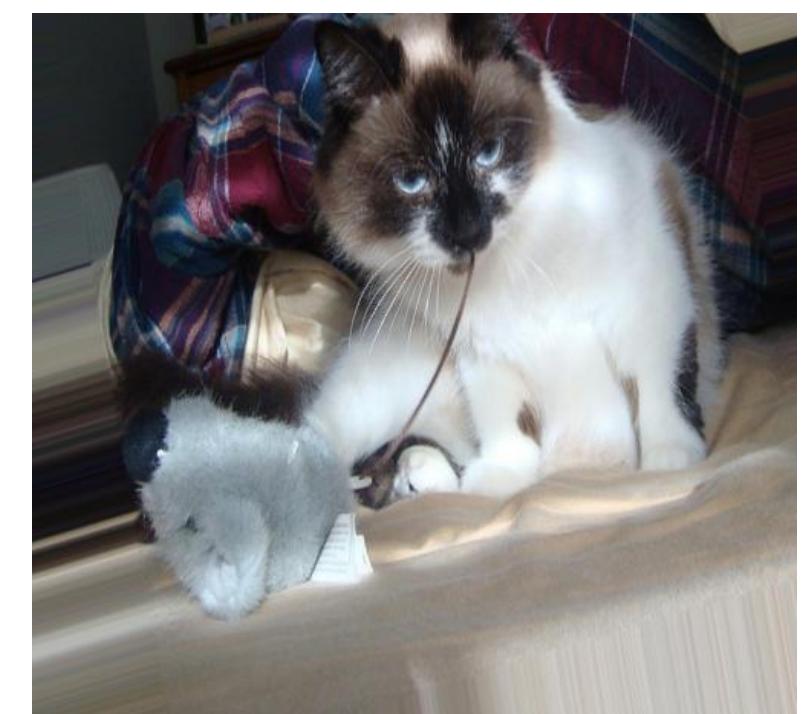
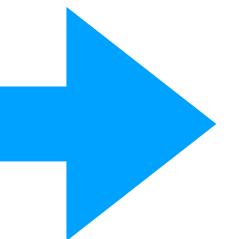
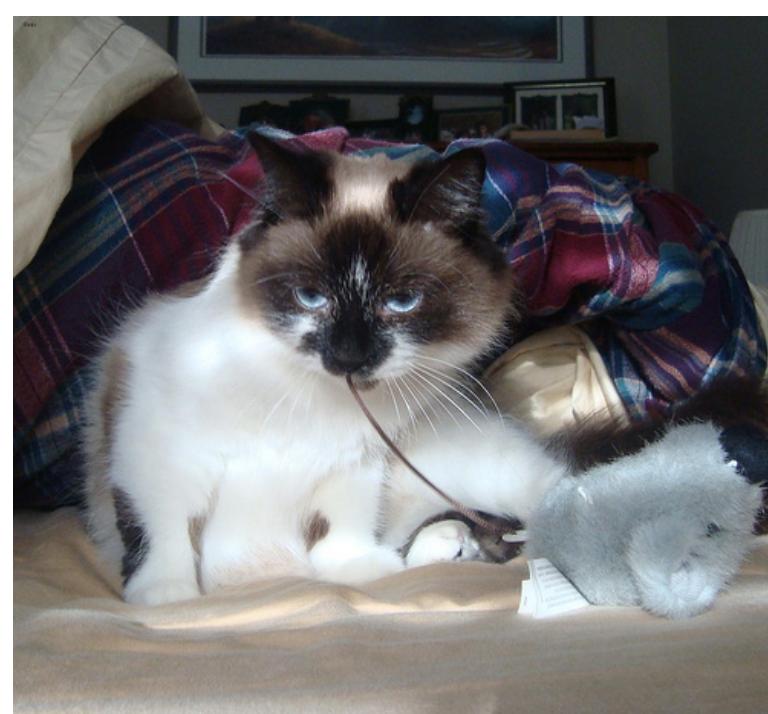
# Training Data



[http://www.robots.ox.ac.uk/  
~vgg/data/pets/](http://www.robots.ox.ac.uk/~vgg/data/pets/)

# Augmenting Data

```
In [ ]: from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img  
  
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```



```
In [69]: from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img

train_data_dir = 'data/train'
validation_data_dir = 'data/validation'
nb_train_samples = 2000
nb_validation_samples = 800
batch_size = 16

# this is the augmentation configuration we will use for training
train_datagen = ImageDataGenerator(
    rescale=1. / 255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)

# this is the augmentation configuration we will use for testing:
# only rescaling
validation_datagen = ImageDataGenerator(rescale=1. / 255)
```

```
Found 2000 images belonging to 2 classes.
Found 800 images belonging to 2 classes.
```

```
In [ ]: train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='binary')

validation_generator = validation_datagen.flow_from_directory(
    validation_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='binary')
```

**Network input size**

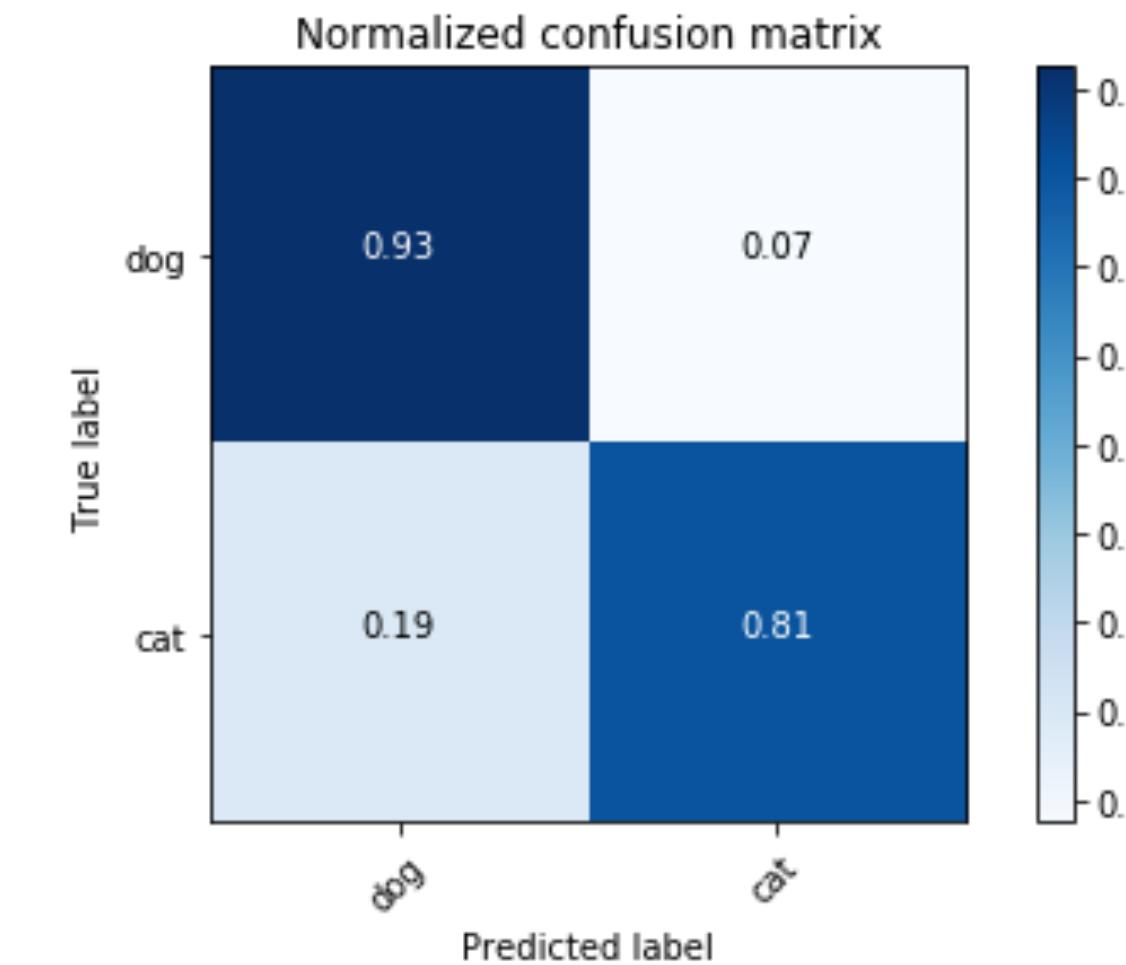
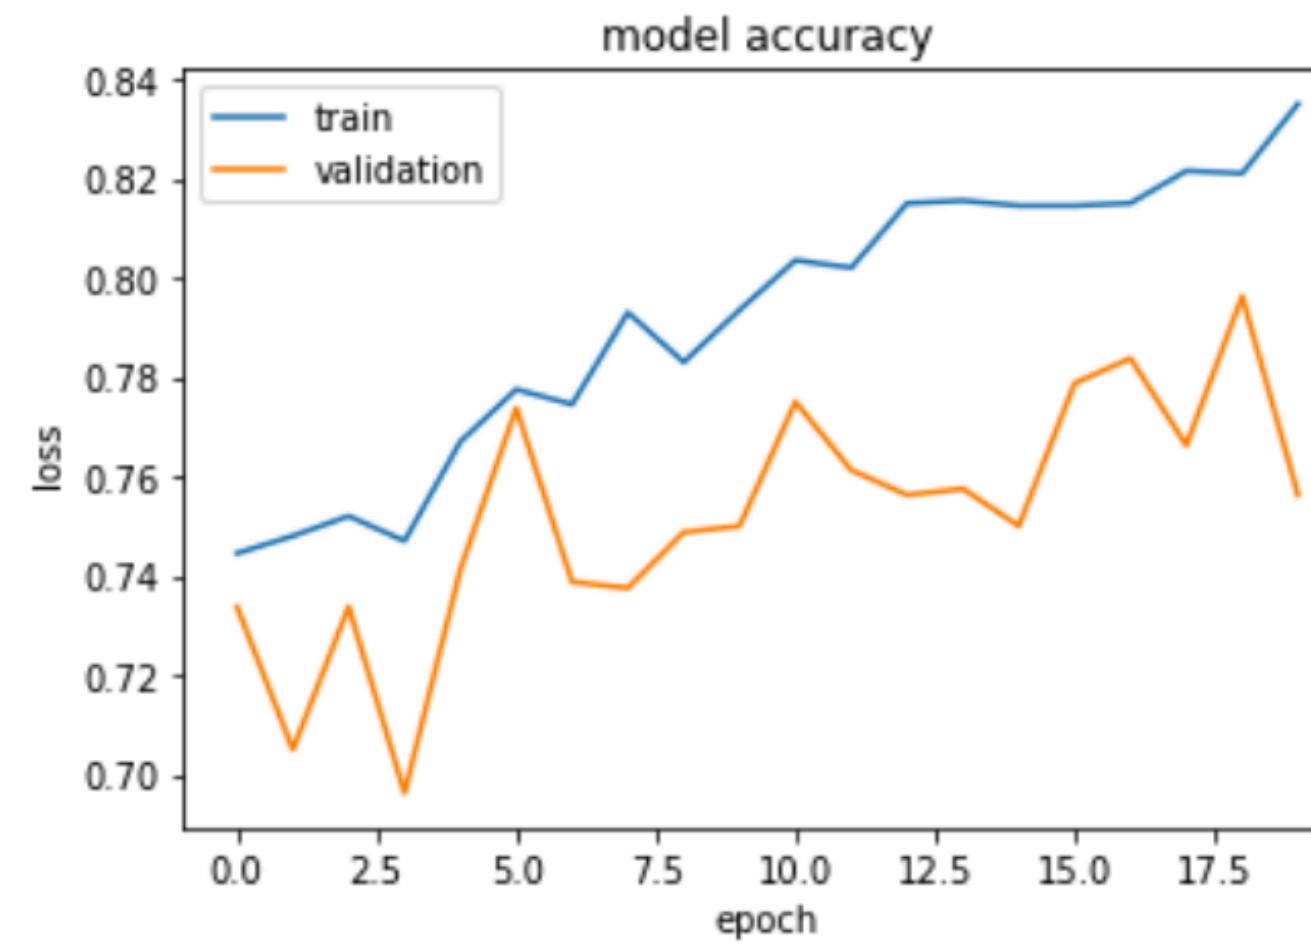
# Training

```
In [20]: history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=nb_train_samples // batch_size,  
    epochs=20,  
    validation_data=validation_generator,  
    validation_steps=nb_validation_samples // batch_size)  
  
model.save_weights('first_try.h5')  
  
Epoch 1/20  
125/125 [=====] - 60s 481ms/step - loss: 0.5441 - acc  
: 0.7445 - val_loss: 0.5389 - val_acc: 0.7338  
Epoch 2/20  
125/125 [=====] - 59s 473ms/step - loss: 0.5246 - acc  
: 0.7480 - val_loss: 0.5549 - val_acc: 0.7050  
Epoch 3/20  
125/125 [=====] - 63s 504ms/step - loss: 0.5101 - acc  
: 0.7520 - val_loss: 0.5298 - val_acc: 0.7338  
Epoch 4/20  
125/125 [=====] - 66s 530ms/step - loss: 0.5276 - acc  
: 0.7470 - val_loss: 0.5853 - val_acc: 0.6963  
Epoch 5/20  
125/125 [=====] - 78s 625ms/step - loss: 0.5004 - acc  
: 0.7670 - val_loss: 0.5165 - val_acc: 0.7412  
Epoch 6/20  
125/125 [=====] - 73s 586ms/step - loss: 0.4802 - acc  
: 0.7775 - val_loss: 0.4903 - val_acc: 0.7738
```

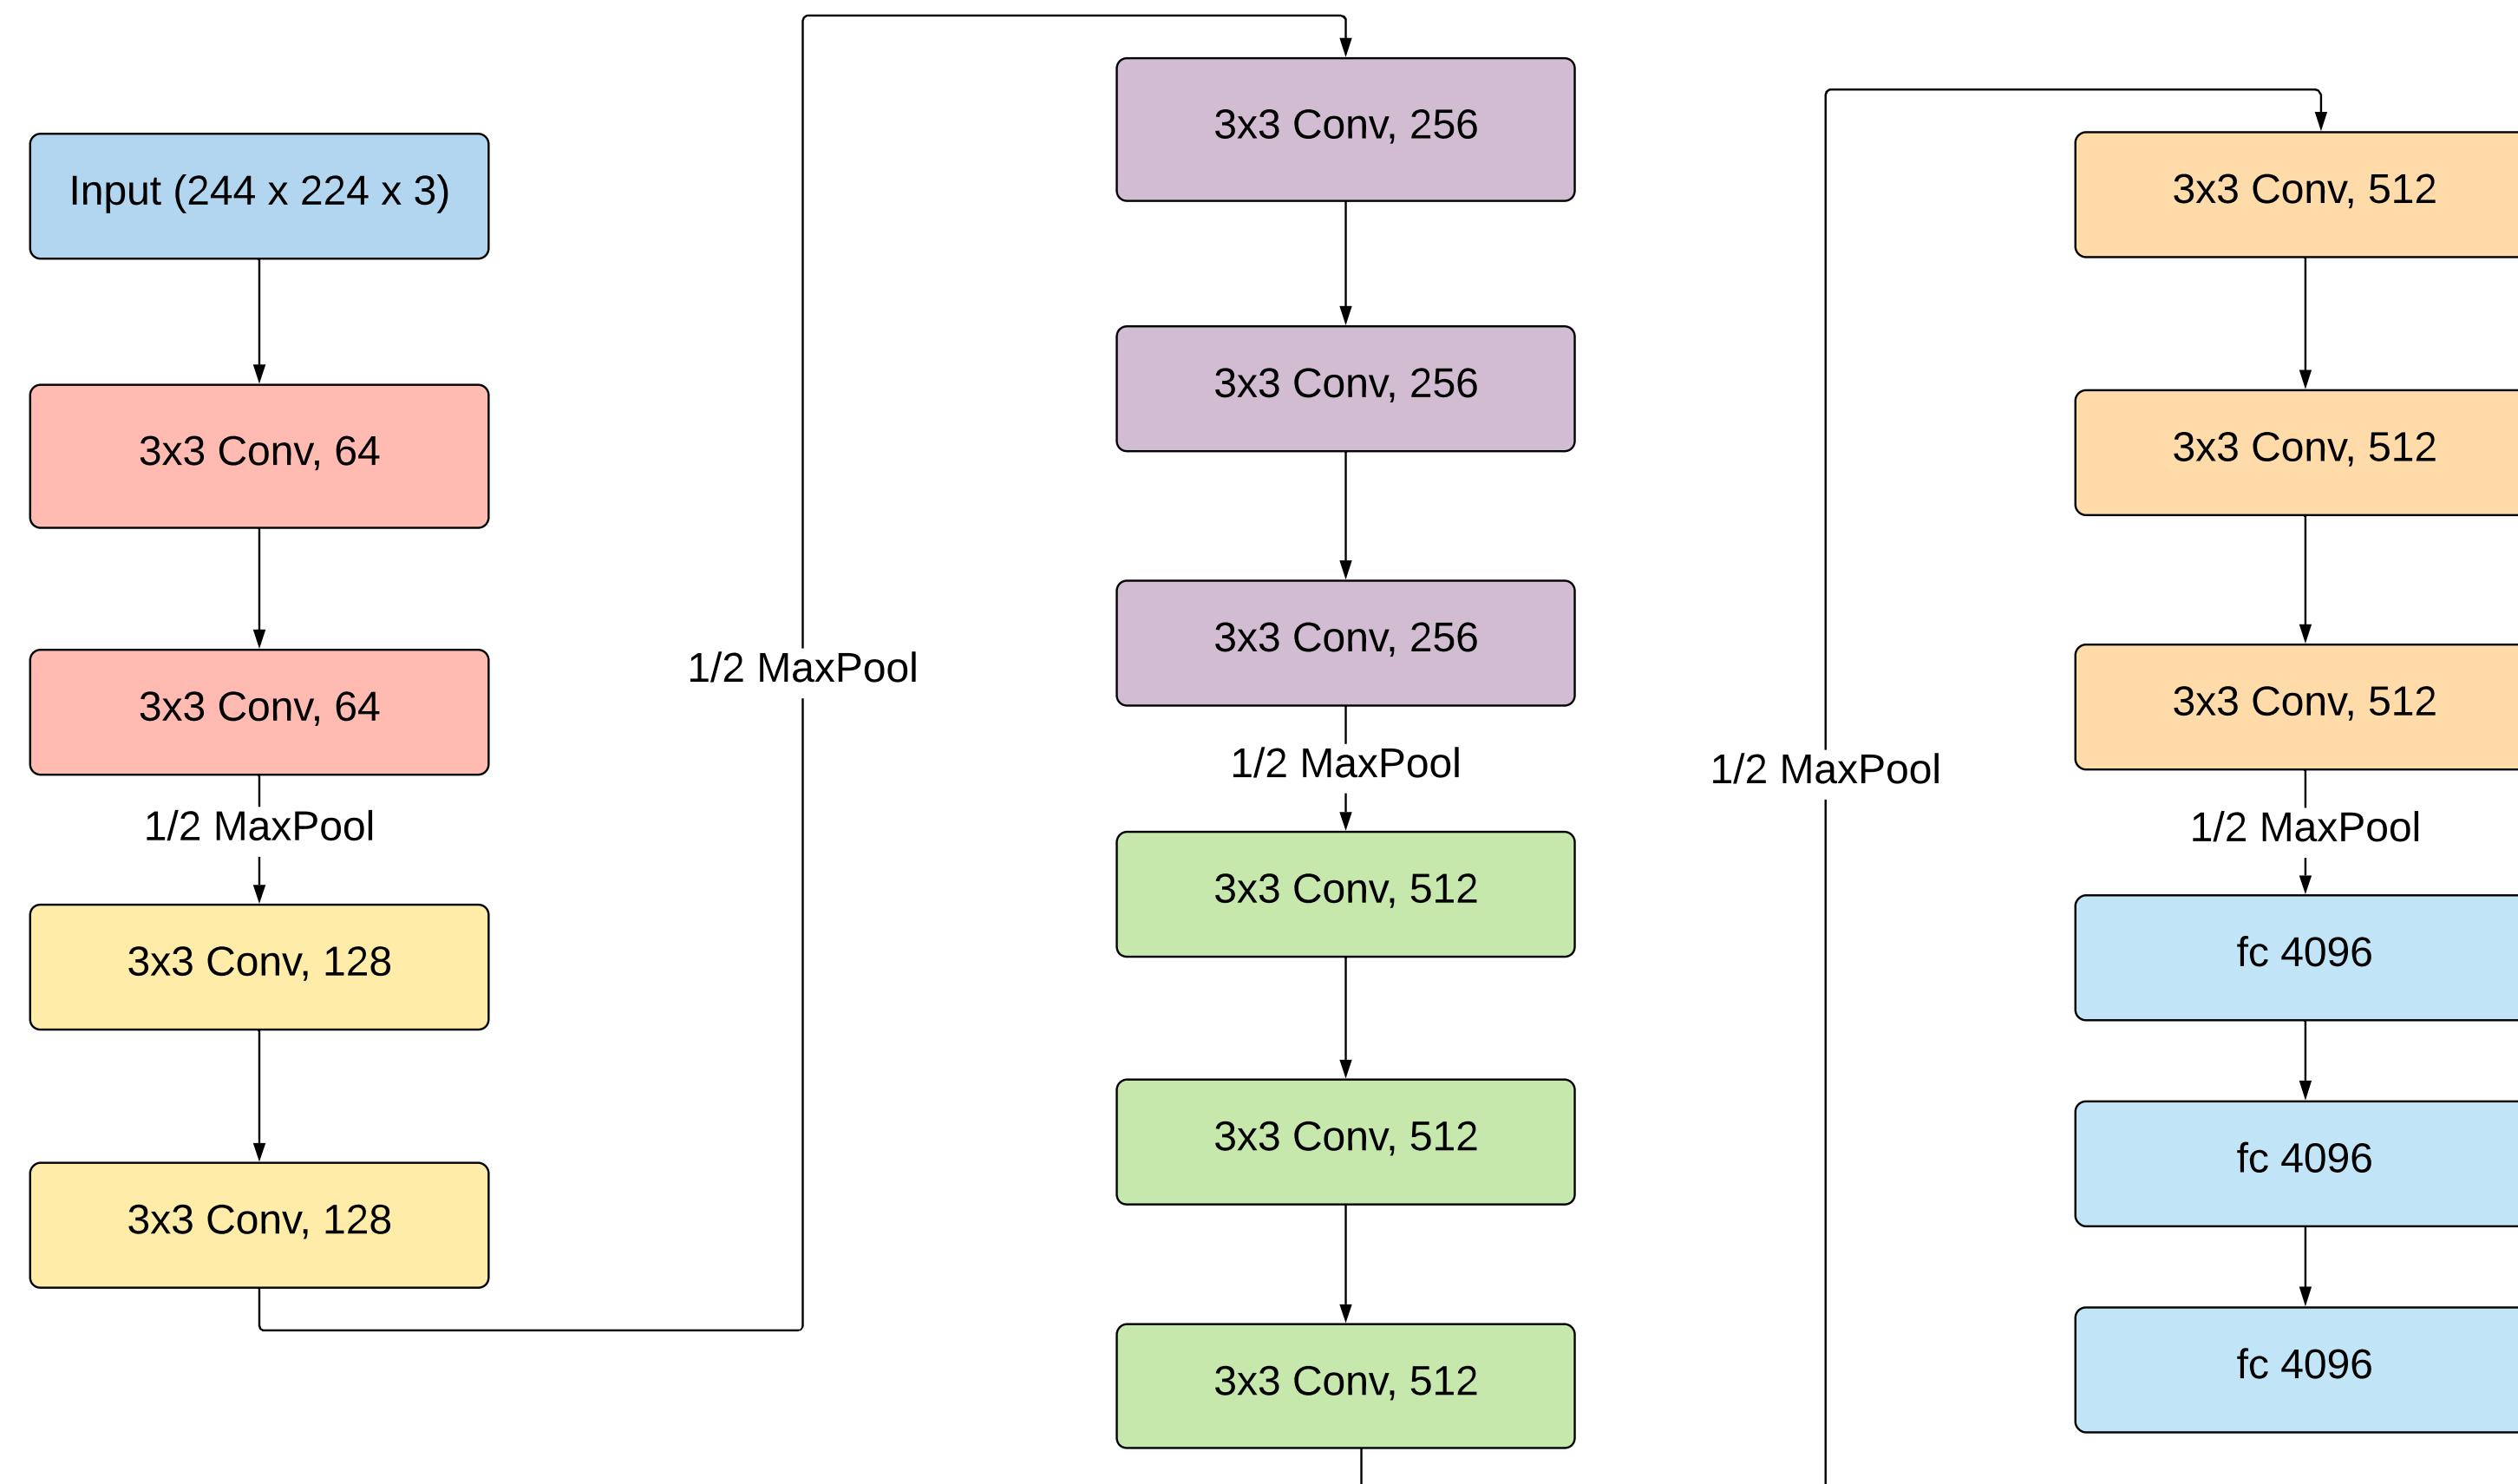
# Evaluation

```
In [22]: import matplotlib.pyplot as plt

plt.plot(history.history['acc'], label='train')
plt.plot(history.history['val_acc'], label = 'validation')
plt.title('model accuracy')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.legend(loc='upper left')
plt.show()
```



# VGG16

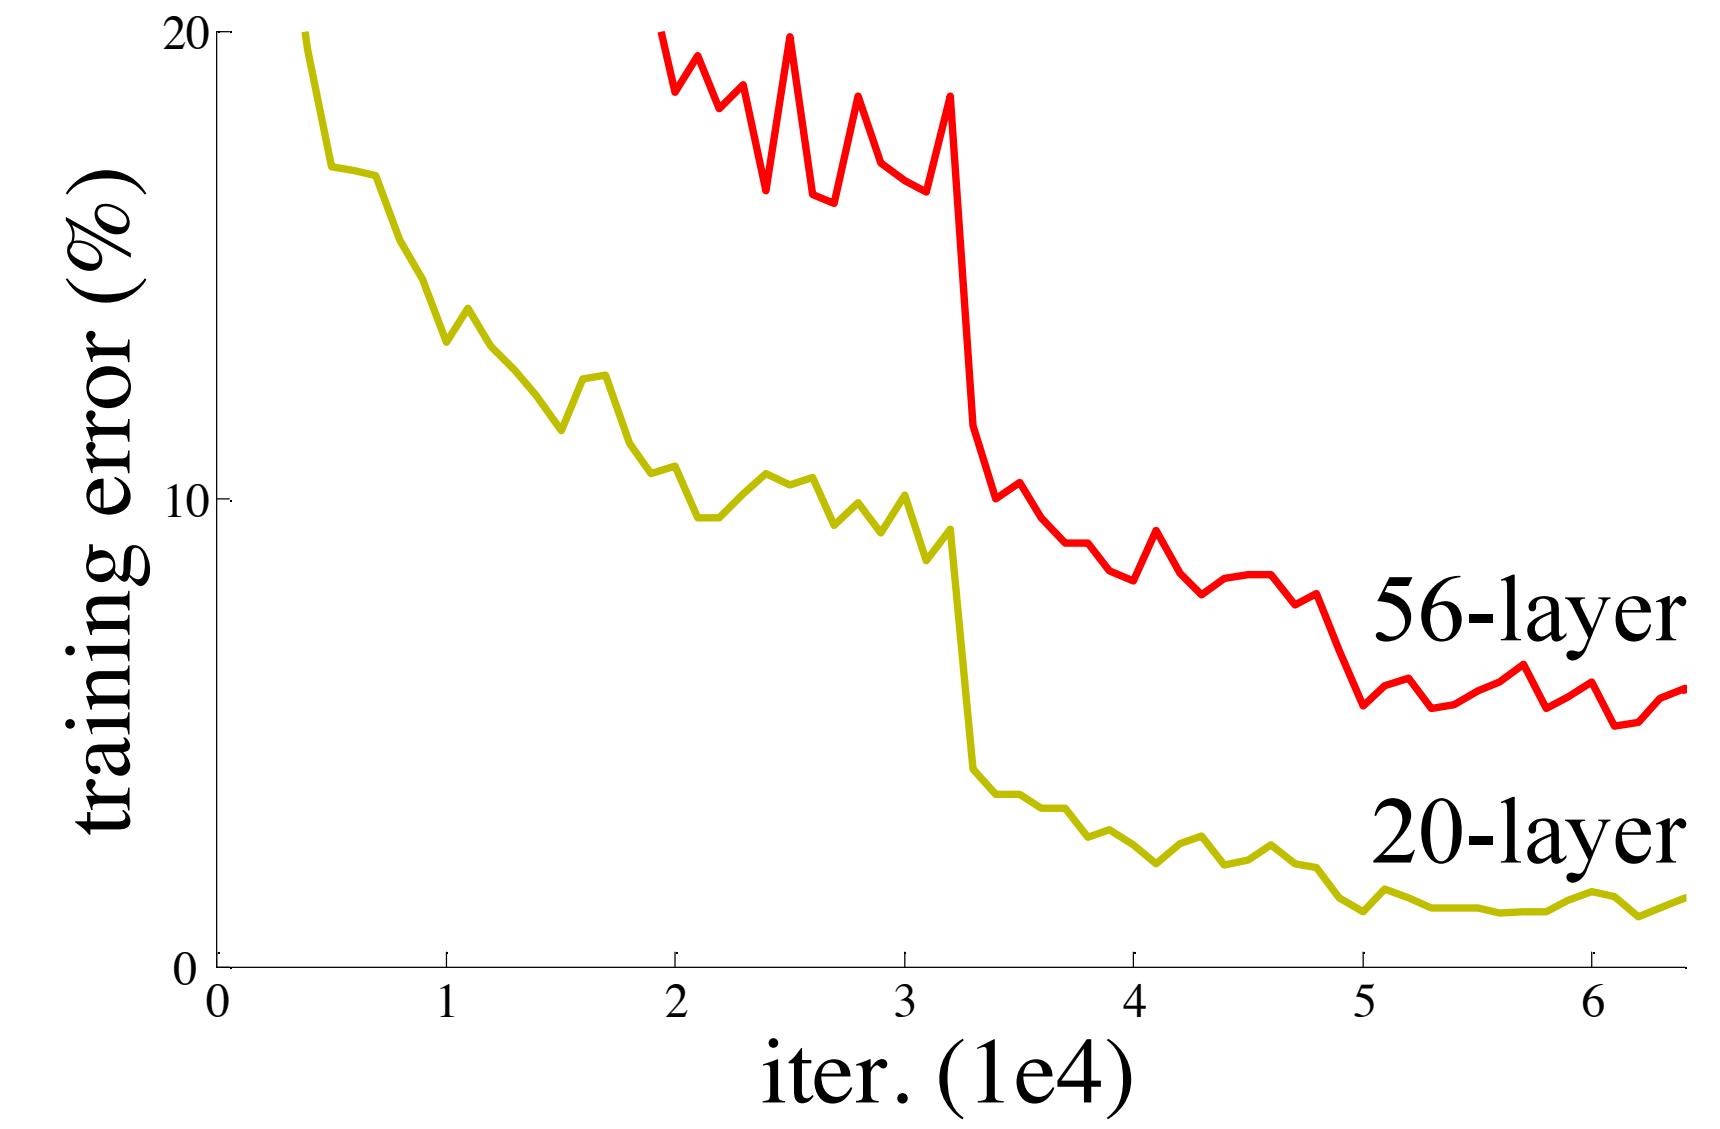


# ResNet

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
				3×3 max pool, stride 2		
conv2_x	56×56	$\begin{bmatrix} 3\times3, 64 \\ 3\times3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 64 \\ 3\times3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 64 \\ 3\times3, 64 \\ 1\times1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3\times3, 128 \\ 3\times3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 128 \\ 3\times3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1\times1, 128 \\ 3\times3, 128 \\ 1\times1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3\times3, 256 \\ 3\times3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 256 \\ 3\times3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1\times1, 256 \\ 3\times3, 256 \\ 1\times1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3\times3, 512 \\ 3\times3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3\times3, 512 \\ 3\times3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 512 \\ 3\times3, 512 \\ 1\times1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1\times1, 512 \\ 3\times3, 512 \\ 1\times1, 2048 \end{bmatrix} \times 3$	
	1×1			average pool, 1000-d fc, softmax		
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

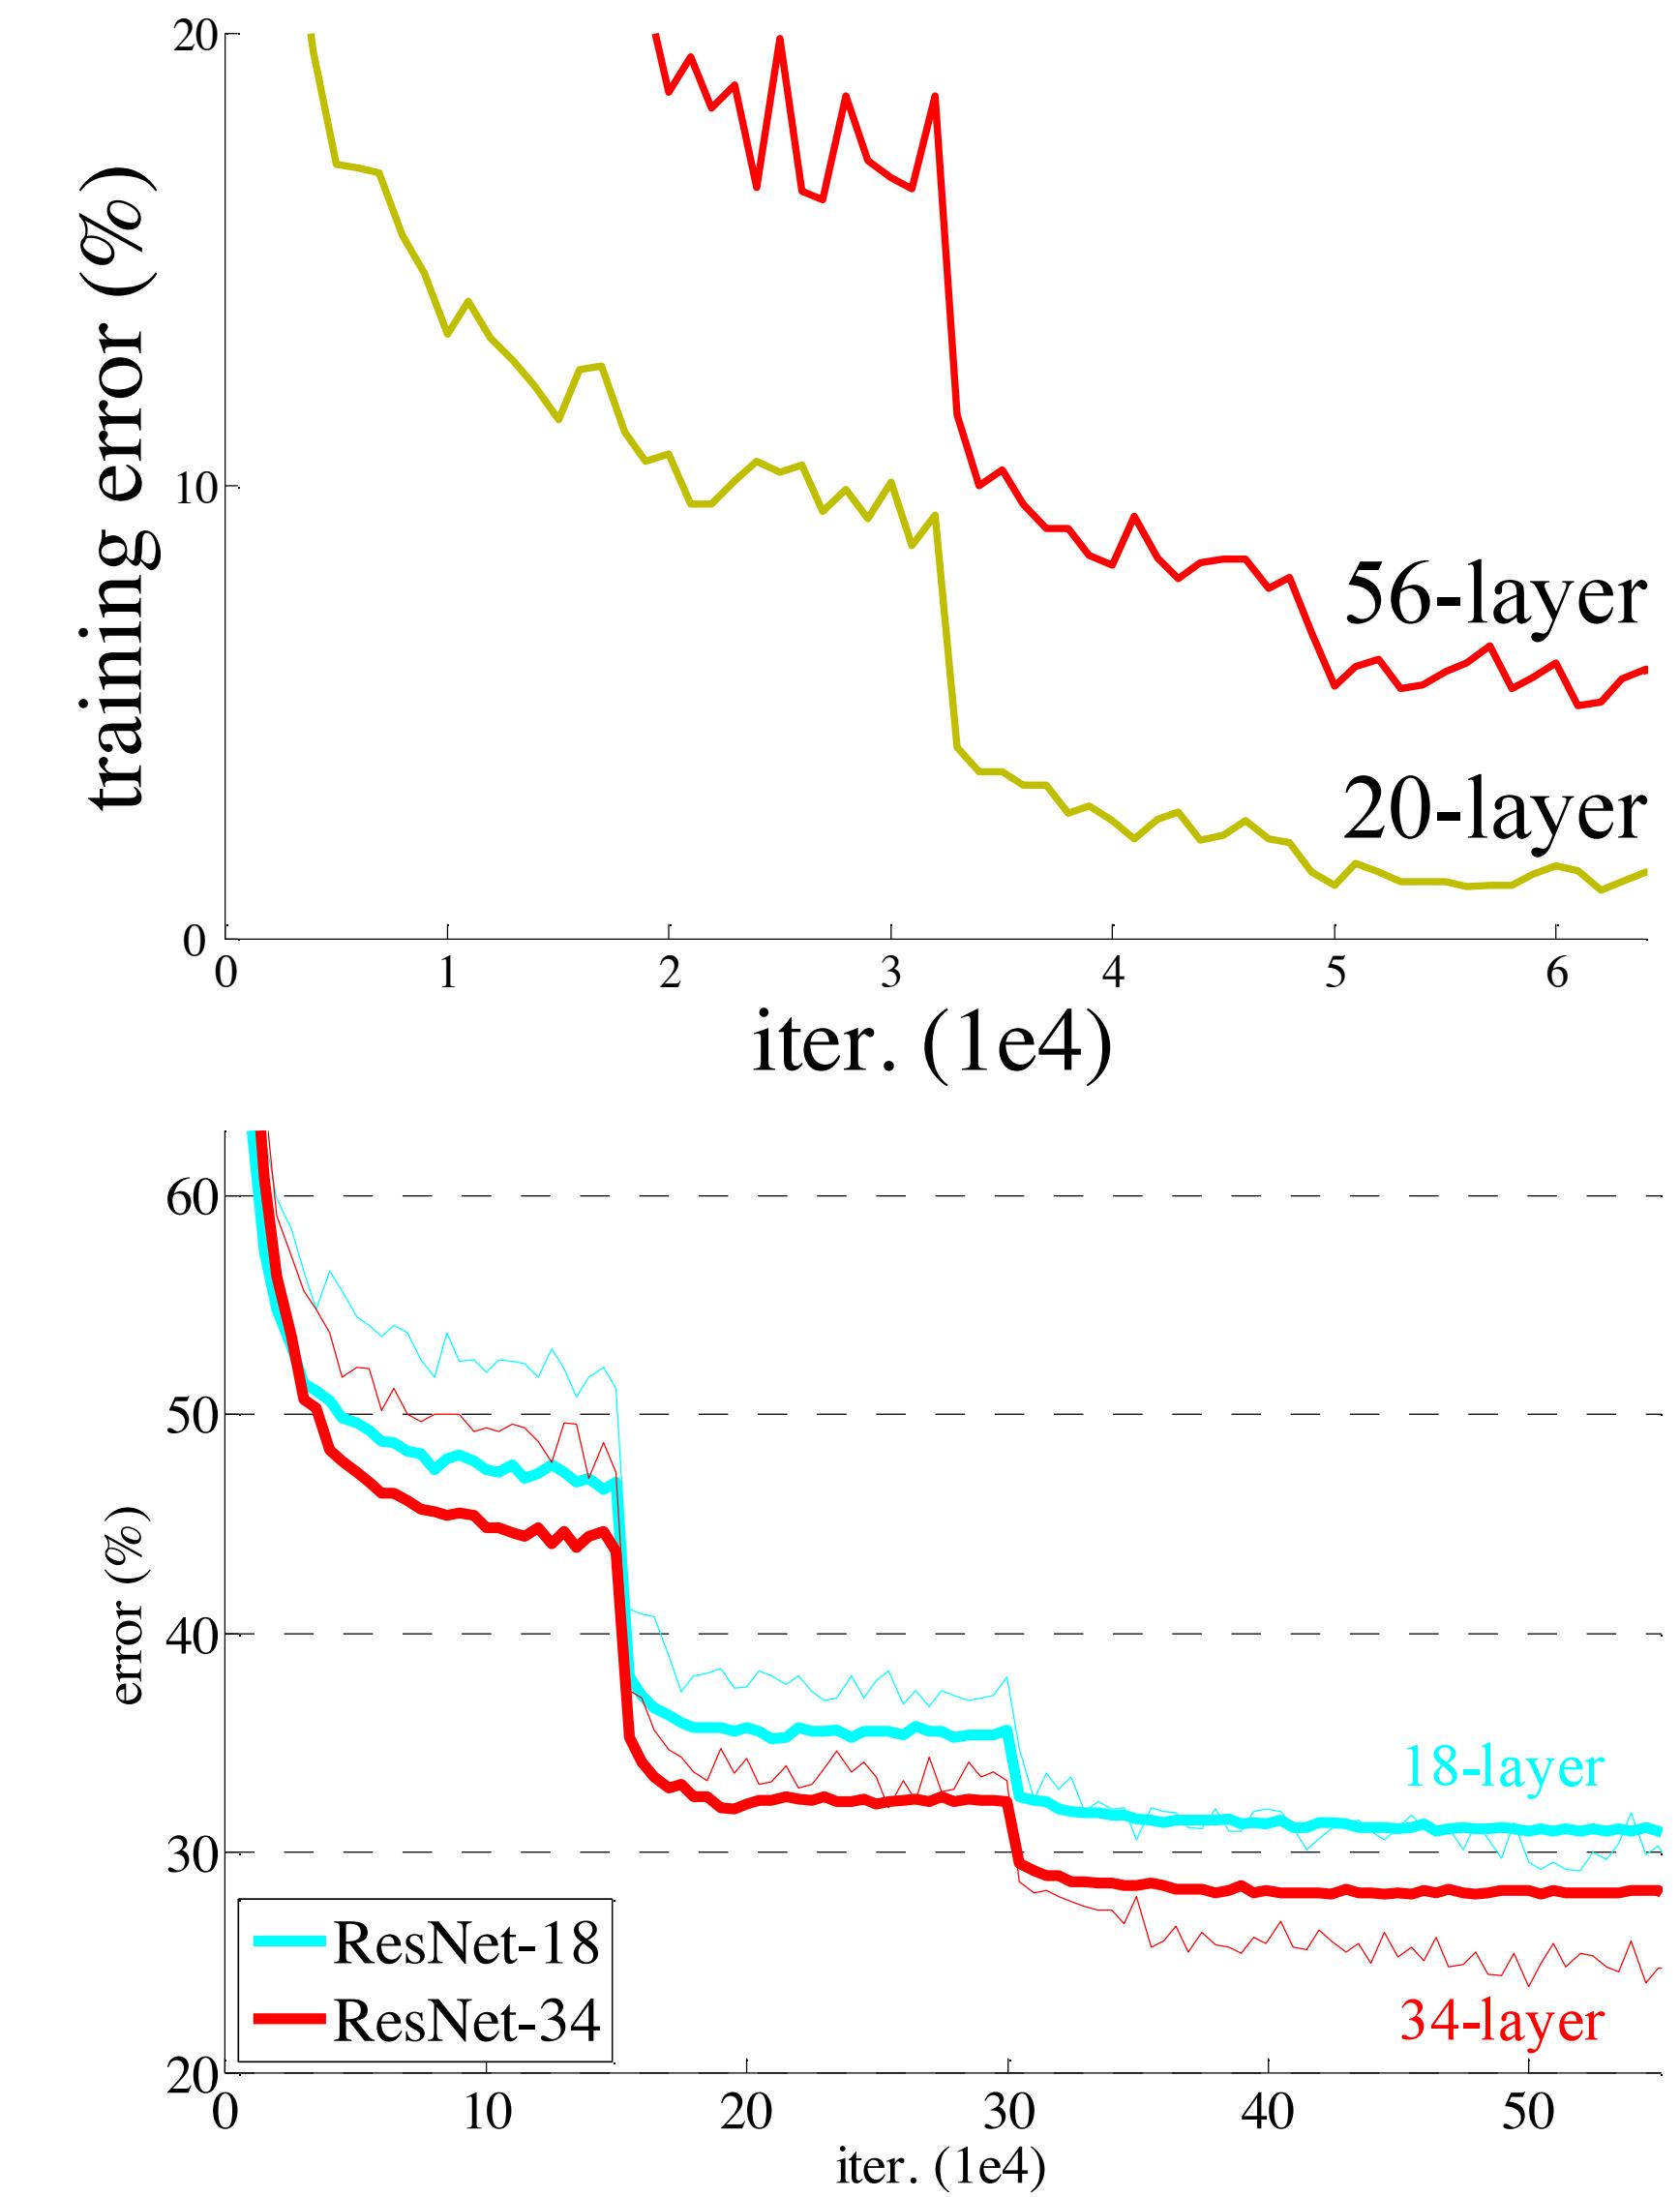
# ResNet

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56	$\left[ \begin{array}{l} 3 \times 3, 64 \\ 3 \times 3, 64 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3 \times 3, 64 \\ 3 \times 3, 64 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{array} \right] \times 3$
conv3_x	28×28	$\left[ \begin{array}{l} 3 \times 3, 128 \\ 3 \times 3, 128 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3 \times 3, 128 \\ 3 \times 3, 128 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{array} \right] \times 8$
conv4_x	14×14	$\left[ \begin{array}{l} 3 \times 3, 256 \\ 3 \times 3, 256 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3 \times 3, 256 \\ 3 \times 3, 256 \end{array} \right] \times 6$	$\left[ \begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 6$	$\left[ \begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 23$	$\left[ \begin{array}{l} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{array} \right] \times 36$
conv5_x	7×7	$\left[ \begin{array}{l} 3 \times 3, 512 \\ 3 \times 3, 512 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3 \times 3, 512 \\ 3 \times 3, 512 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{array} \right] \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

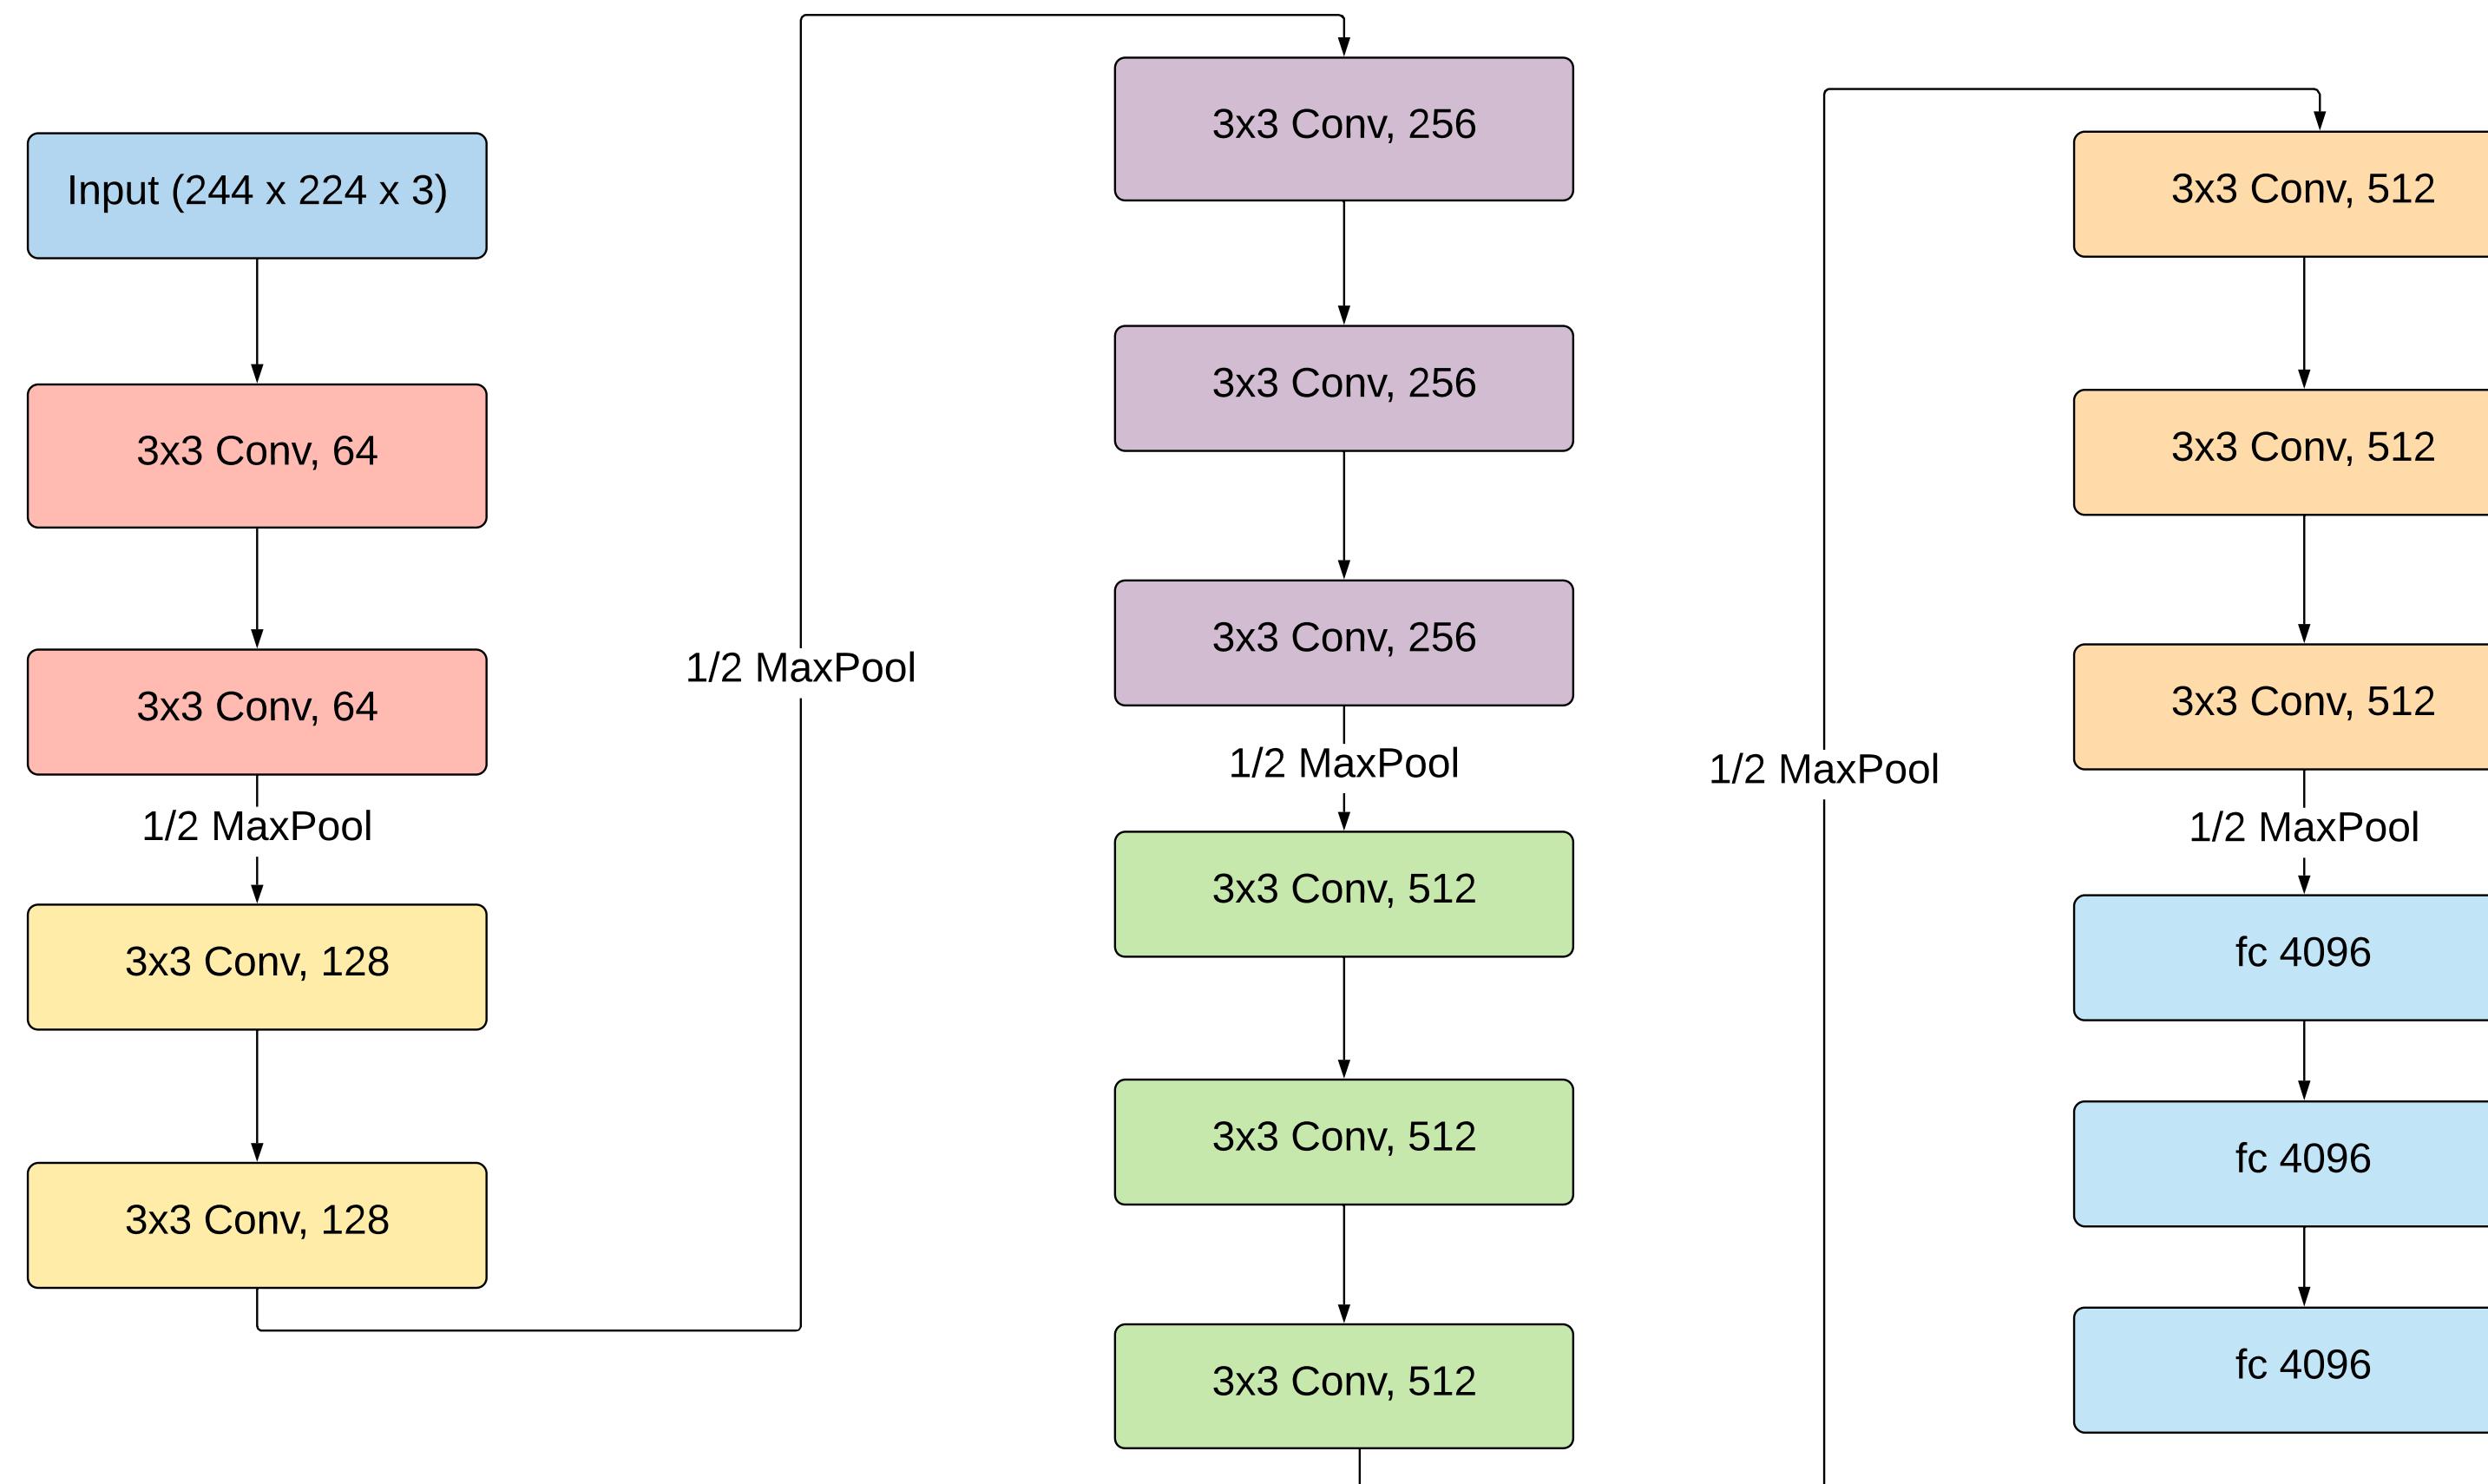


# ResNet

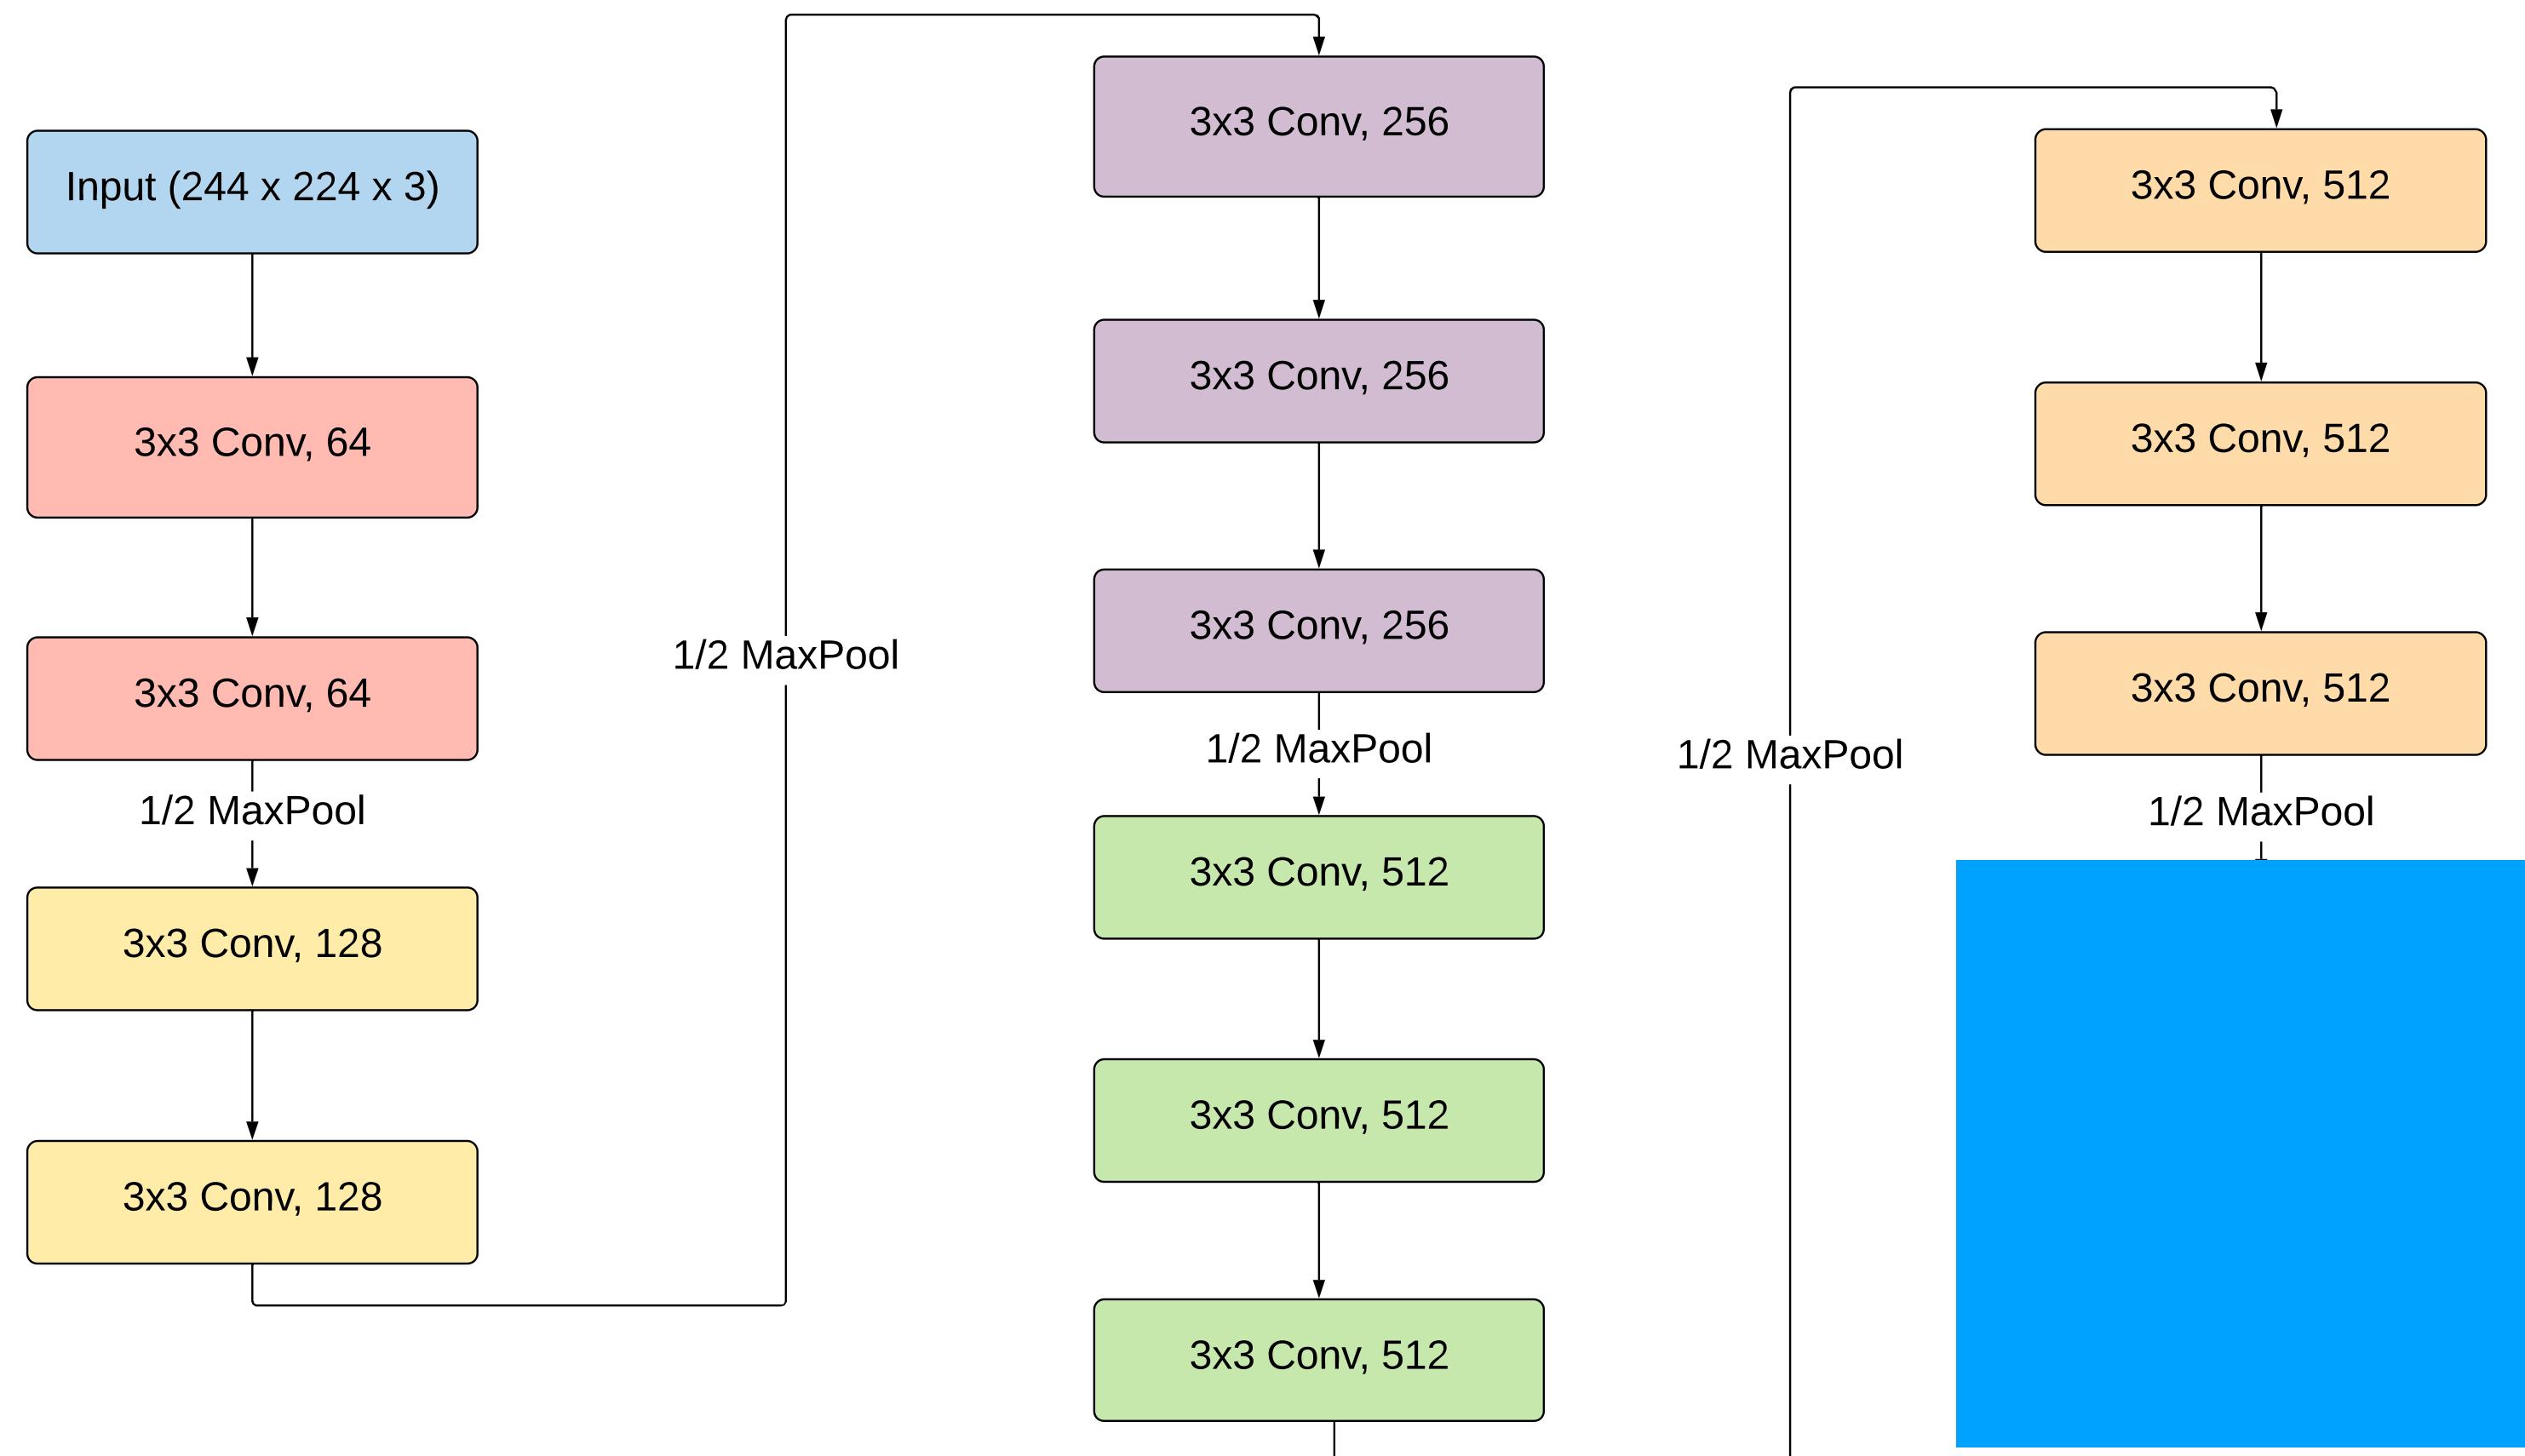
layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56	$\left[ \begin{array}{l} 3\times 3, 64 \\ 3\times 3, 64 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3\times 3, 64 \\ 3\times 3, 64 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{array} \right] \times 3$
conv3_x	28×28	$\left[ \begin{array}{l} 3\times 3, 128 \\ 3\times 3, 128 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3\times 3, 128 \\ 3\times 3, 128 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{array} \right] \times 4$	$\left[ \begin{array}{l} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{array} \right] \times 8$
conv4_x	14×14	$\left[ \begin{array}{l} 3\times 3, 256 \\ 3\times 3, 256 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3\times 3, 256 \\ 3\times 3, 256 \end{array} \right] \times 6$	$\left[ \begin{array}{l} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{array} \right] \times 6$	$\left[ \begin{array}{l} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{array} \right] \times 23$	$\left[ \begin{array}{l} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{array} \right] \times 36$
conv5_x	7×7	$\left[ \begin{array}{l} 3\times 3, 512 \\ 3\times 3, 512 \end{array} \right] \times 2$	$\left[ \begin{array}{l} 3\times 3, 512 \\ 3\times 3, 512 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{array} \right] \times 3$	$\left[ \begin{array}{l} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{array} \right] \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$



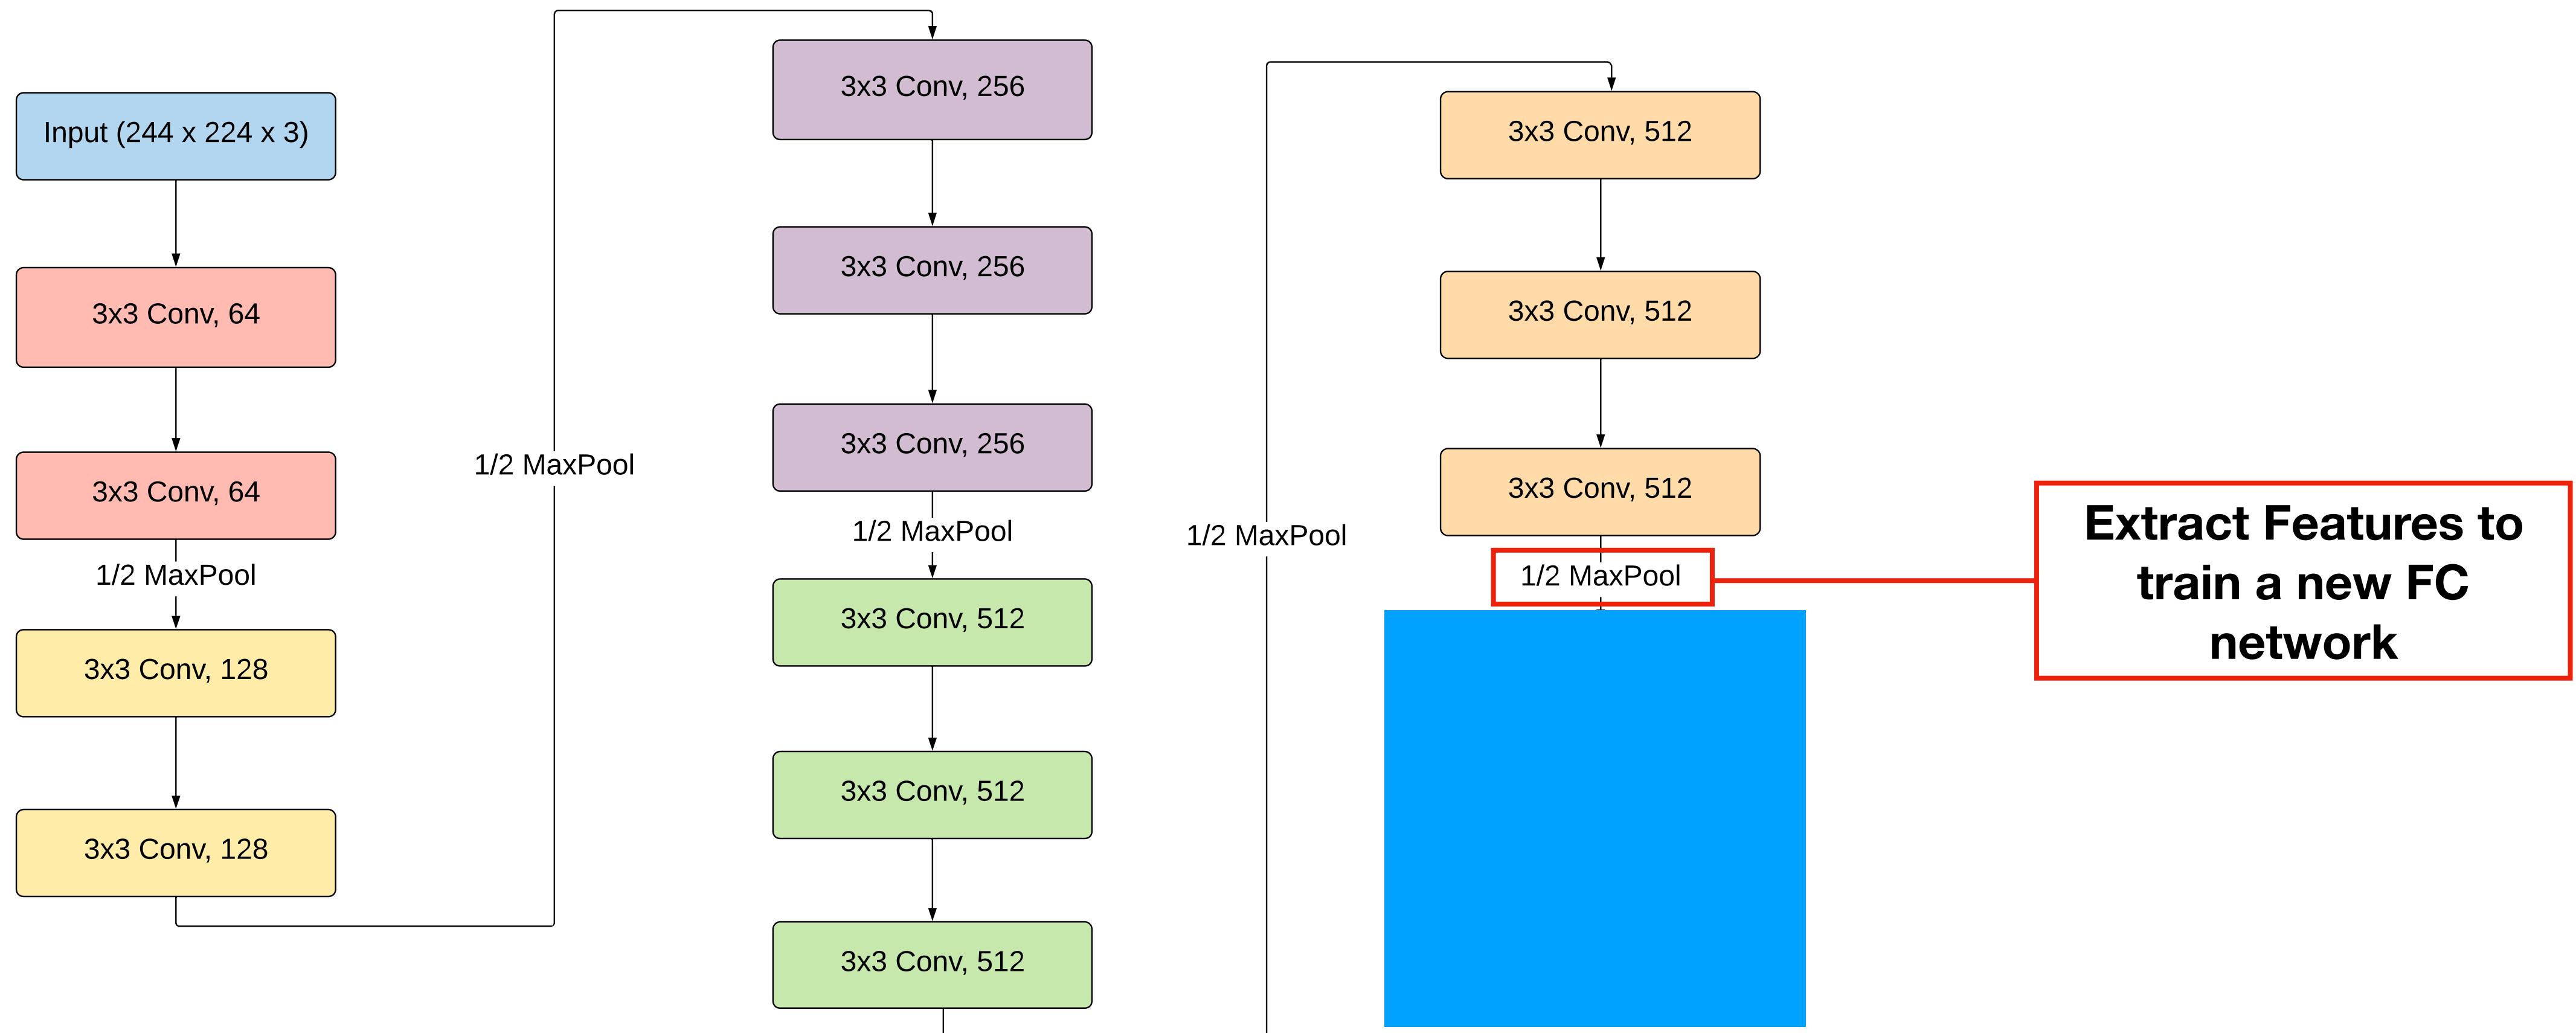
# Bottleneck Features



# Bottleneck Features



# Bottleneck Features



# Bottleneck Features

```
In [ ]: import keras
from keras import applications
from keras import models
from keras.layers import Dropout, Flatten, Dense
import numpy as np

batch_size = 16

# Load the VGG16 network with the pretrained imagenet weights. Don't include the fully connected layers.

model = applications.VGG16(include_top=False, weights='imagenet')

# Generate feature map for training data
generator = datagen.flow_from_directory(
    'data/train',
    target_size=(150, 150),
    batch_size=batch_size,
    class_mode=None,
    shuffle=False)

bottleneck_features_train = model.predict_generator(generator, 2000//batch_size)

# save the output as a Numpy array
np.save('bottleneck_features_train.npy', bottleneck_features_train)
```

**Load VGG16 model with pertained weights and no FC layers**

**Generate bottleneck features**

# Train FC layer

```
In [29]: from keras.utils.np_utils import to_categorical
datagen_top = ImageDataGenerator(rescale=1./255)
generator_top = datagen_top.flow_from_directory(
    train_data_dir,
    target_size=(img_width, img_height),
    batch_size=batch_size,
    class_mode='categorical',
    shuffle=False)

num_classes = len(generator_top.class_indices)

# load the bottleneck features saved earlier
train_data = np.load('bottleneck_features_train.npy')

# get the class labels for the training data, in the original order
train_labels = generator_top.classes
```

Load bottleneck  
features for training

Load labels

# Train FC layer

```
In [ ]: model = Sequential()
model.add(Flatten(input_shape=train_data.shape[1:]))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=[ 'accuracy'])

history = model.fit(train_data, train_labels,
                     epochs=50,
                     batch_size=batch_size,
                     validation_data=(validation_data, validation_labels))
model.save_weights('bottleneck_fc_model.h5')
```

# Train FC layer

```
In [ ]: model = Sequential()
model.add(Flatten(input_shape=train_data.shape[1:]))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=[ 'accuracy'])

history = model.fit(train_data, train_labels,
                     epochs=50,
                     batch_size=batch_size,
                     validation_data=(validation_data, validation_labels))
model.save_weights('bottleneck_fc_model.h5')
```

Pass bottleneck  
features and labels to  
fitting

# Train FC layer

```
In [ ]: model = Sequential()
model.add(Flatten(input_shape=train_data.shape[1:]))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=[ 'accuracy'])

history = model.fit(train_data, train_labels,
                     epochs=50,
                     batch_size=batch_size,
                     validation_data=(validation_data, validation_labels))
model.save_weights('bottleneck_fc_model.h5')
```

Pass bottleneck  
features and labels to  
fitting

Save FC weights

# Fine Tune Network

```
In [65]: # build model
base_model = applications.VGG16(input_shape=input_shape, include_top=False, weight
s='imagenet')

# Create a model
fullyconnected_model = Sequential()
fullyconnected_model.add(Flatten(input_shape=base_model.output_shape[1:]))
fullyconnected_model.add(Dense(256, activation='relu'))
fullyconnected_model.add(Dropout(0.5))
fullyconnected_model.add(Dense(1, activation='sigmoid'))

fullyconnected_model.load_weights('bottleneck_fc_model.h5')

model = models.Model(inputs= base_model.input, outputs= fullyconnected_model(base_
model.output))

for layer in model.layers[:-3]:
    layer.trainable = False

adam=keras.optimizers.Adam(lr=0.0001)
model.compile(optimizer=adam,
              loss='binary_crossentropy',
              metrics=['accuracy'])

print('model compiled')

model.summary()
```

# Fine Tune Network

```
In [65]: # build model
base_model = applications.VGG16(input_shape=input_shape, include_top=False, weight
s='imagenet')

# Create a model
fullyconnected_model = Sequential()
fullyconnected_model.add(Flatten(input_shape=base_model.output_shape[1:]))
fullyconnected_model.add(Dense(256, activation='relu'))
fullyconnected_model.add(Dropout(0.5))
fullyconnected_model.add(Dense(1, activation='sigmoid'))

fullyconnected_model.load_weights('bottleneck_fc_model.h5')
```

model = models.Model(inputs= base\_model.input, outputs= fullyconnected\_model(base\_
model.output))

**for** layer **in** model.layers[:-3]:
 layer.trainable = **False**

adam=keras.optimizers.Adam(lr=0.0001)
model.compile(optimizer=adam,
 loss='binary\_crossentropy',
 metrics=['accuracy'])

print('model compiled')

model.summary()

Load FC weights

# Fine Tune Network

```
In [65]: # build model
base_model = applications.VGG16(input_shape=input_shape, include_top=False, weight
s='imagenet')

# Create a model
fullyconnected_model = Sequential()
fullyconnected_model.add(Flatten(input_shape=base_model.output_shape[1:]))
fullyconnected_model.add(Dense(256, activation='relu'))
fullyconnected_model.add(Dropout(0.5))
fullyconnected_model.add(Dense(1, activation='sigmoid'))

fullyconnected_model.load_weights('bottleneck_fc_model.h5')

model = models.Model(inputs= base_model.input, outputs= fullyconnected_model(base_
model.output))

for layer in model.layers[:-3]:
    layer.trainable = False

adam=keras.optimizers.Adam(lr=0.0001)
model.compile(optimizer=adam,
              loss='binary_crossentropy',
              metrics=['accuracy'])

print('model compiled')

model.summary()
```

Load FC weights

Combine VGG16  
base with the FC  
network

# Fine Tune Network

```
In [65]: # build model
base_model = applications.VGG16(input_shape=input_shape, include_top=False, weights='imagenet')

# Create a model
fullyconnected_model = Sequential()
fullyconnected_model.add(Flatten(input_shape=base_model.output_shape[1:]))
fullyconnected_model.add(Dense(256, activation='relu'))
fullyconnected_model.add(Dropout(0.5))
fullyconnected_model.add(Dense(1, activation='sigmoid'))

fullyconnected_model.load_weights('bottleneck_fc_model.h5')

model = models.Model(inputs= base_model.input, outputs= fullyconnected_model(base_model.output))

for layer in model.layers[:-3]:
    layer.trainable = False

adam=keras.optimizers.Adam(lr=0.0001)
model.compile(optimizer=adam,
              loss='binary_crossentropy',
              metrics=['accuracy'])

print('model compiled')

model.summary()
```

fullyconnected\_model.load\_weights('bottleneck\_fc\_model.h5')

**Load FC weights**

model = models.Model(inputs= base\_model.input, outputs= fullyconnected\_model(base\_model.output))

**Combine VGG16 base with the FC network**

for layer in model.layers[:-3]:  
 layer.trainable = False

**Chose which layers to freeze**

adam=keras.optimizers.Adam(lr=0.0001)  
model.compile(optimizer=adam,  
 loss='binary\_crossentropy',  
 metrics=['accuracy'])

print('model compiled')

model.summary()

Layer (type)	Output Shape	Param #
<hr/>		
input_11 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
sequential_13 (Sequential)	(None, 1)	2097665
<hr/>		
Total params: 16,812,353		
Trainable params: 4,457,473		
Non-trainable params: 12,354,880		

## From Scratch

	precision	recall	f1-score	support
dog	0.84	0.63	0.72	400
cat	0.71	0.88	0.78	400
avg / total	0.77	0.76	0.75	800

## Fine Tuning VGG16

Found 800 images belonging to 2 classes.

	precision	recall	f1-score	support
dog	0.93	0.93	0.93	400
cat	0.93	0.94	0.93	400
avg / total	0.93	0.93	0.93	800

## Using VGG16 Bottleneck Features

	precision	recall	f1-score	support
dog	0.88	0.87	0.87	400
cat	0.87	0.88	0.88	400
avg / total	0.88	0.88	0.87	800